

- **Think** till **left fork** is available; when it is → **Pick it up**
- **Think** till **right fork** is available; when it is → **Pick it up**
- **Eat** for a fixed amount of time when **both forks** are held
- **Put the right fork down** → **Put the left fork down**
- **Repeat** the process


DEADLOCK

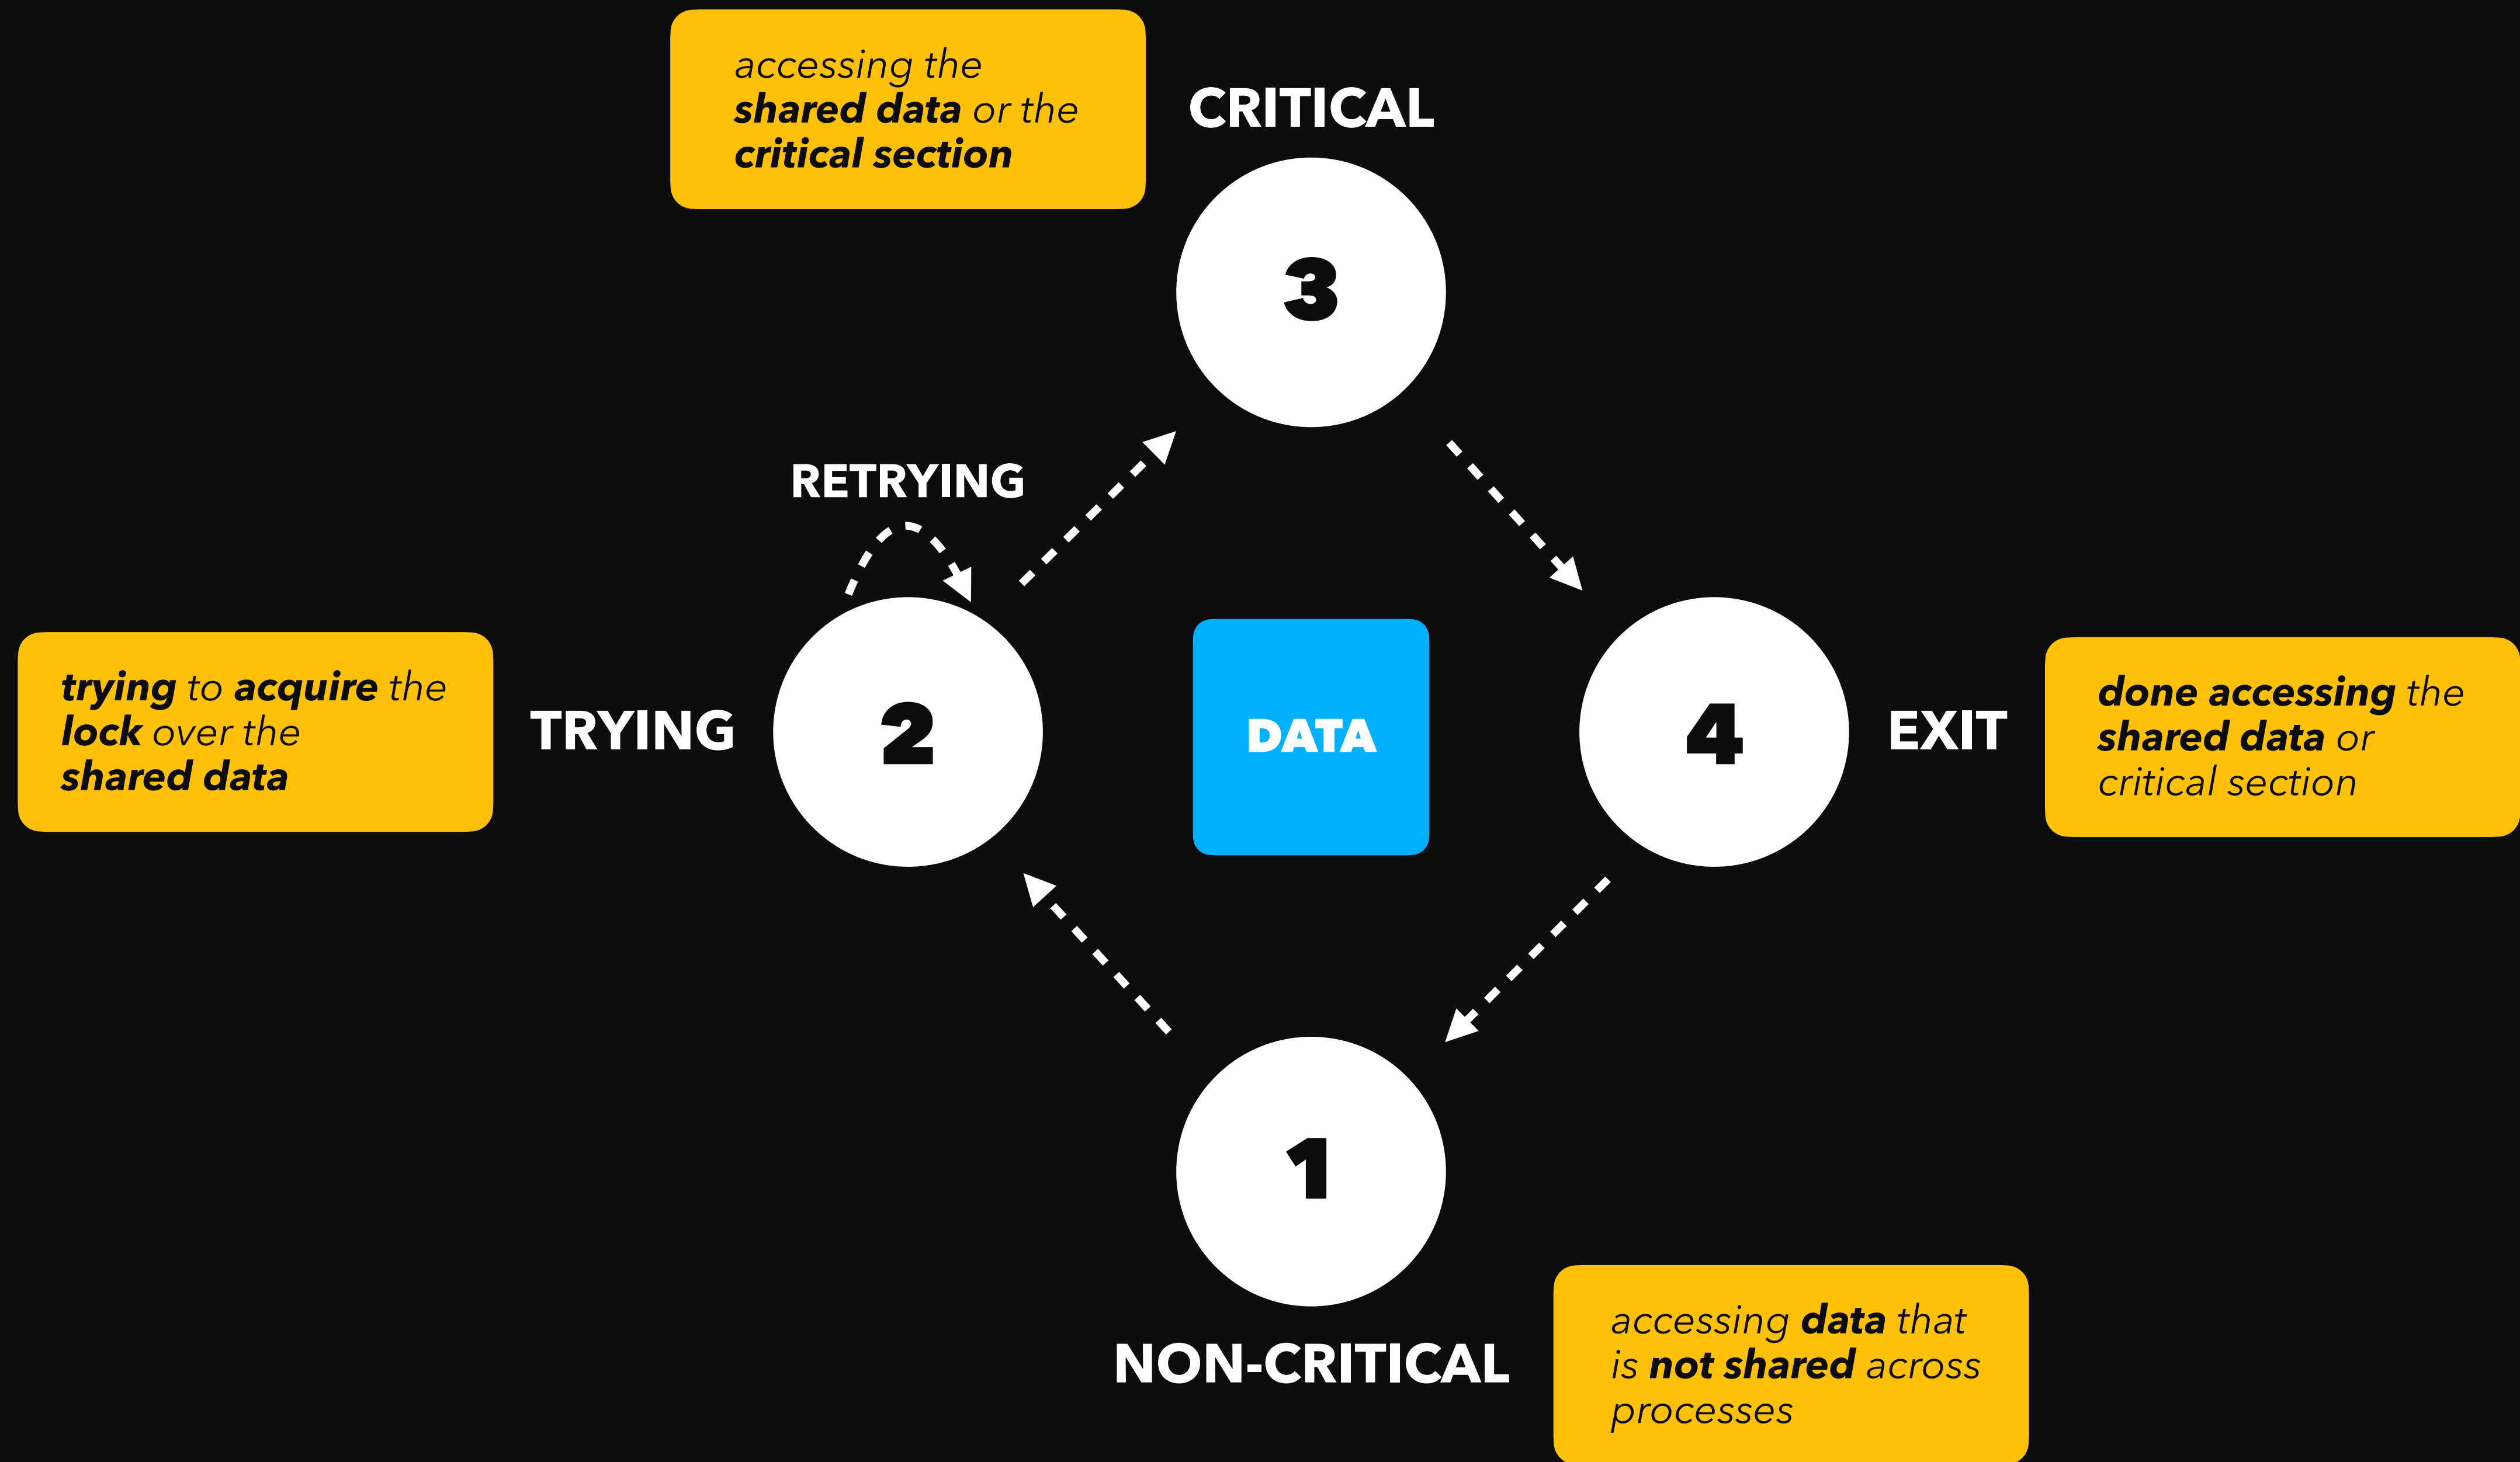
STARVATION

LIVELOCK

MUTUAL EXCLUSION

In computer science **Mutual Exclusion** is a property of **Concurrency Control**, which is instituted for the purpose of preventing **Race Conditions**.





LOCKS	MUTEX
READERS-WRITER LOCKS	RWMUTEX
RECURSIVE LOCKS	UNAVAILABLE
SEMAPHORES	INTERNAL
MONITORS	INTERNAL



1




2



3





1 0 2



ORDER

RESULT

CORRECTNESS



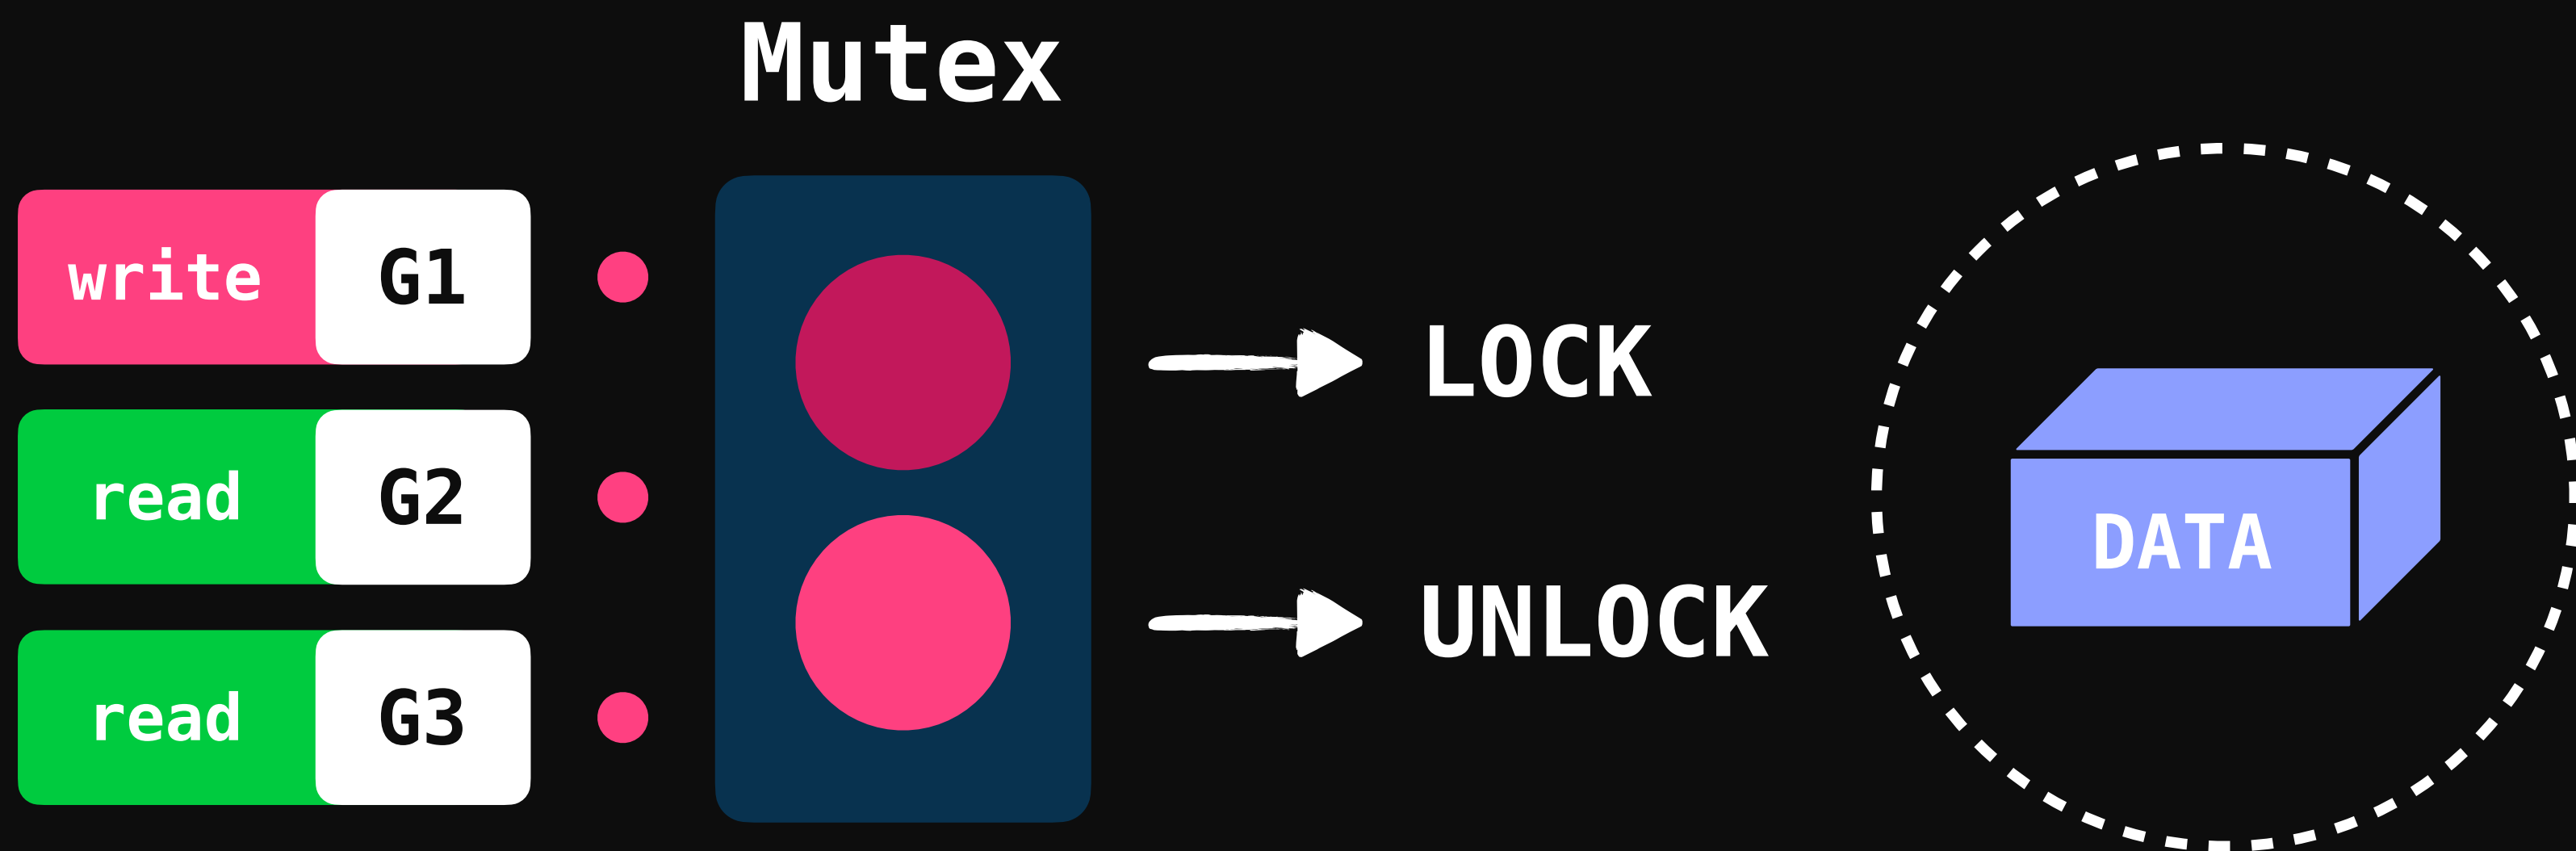


READ WRITE ONCE

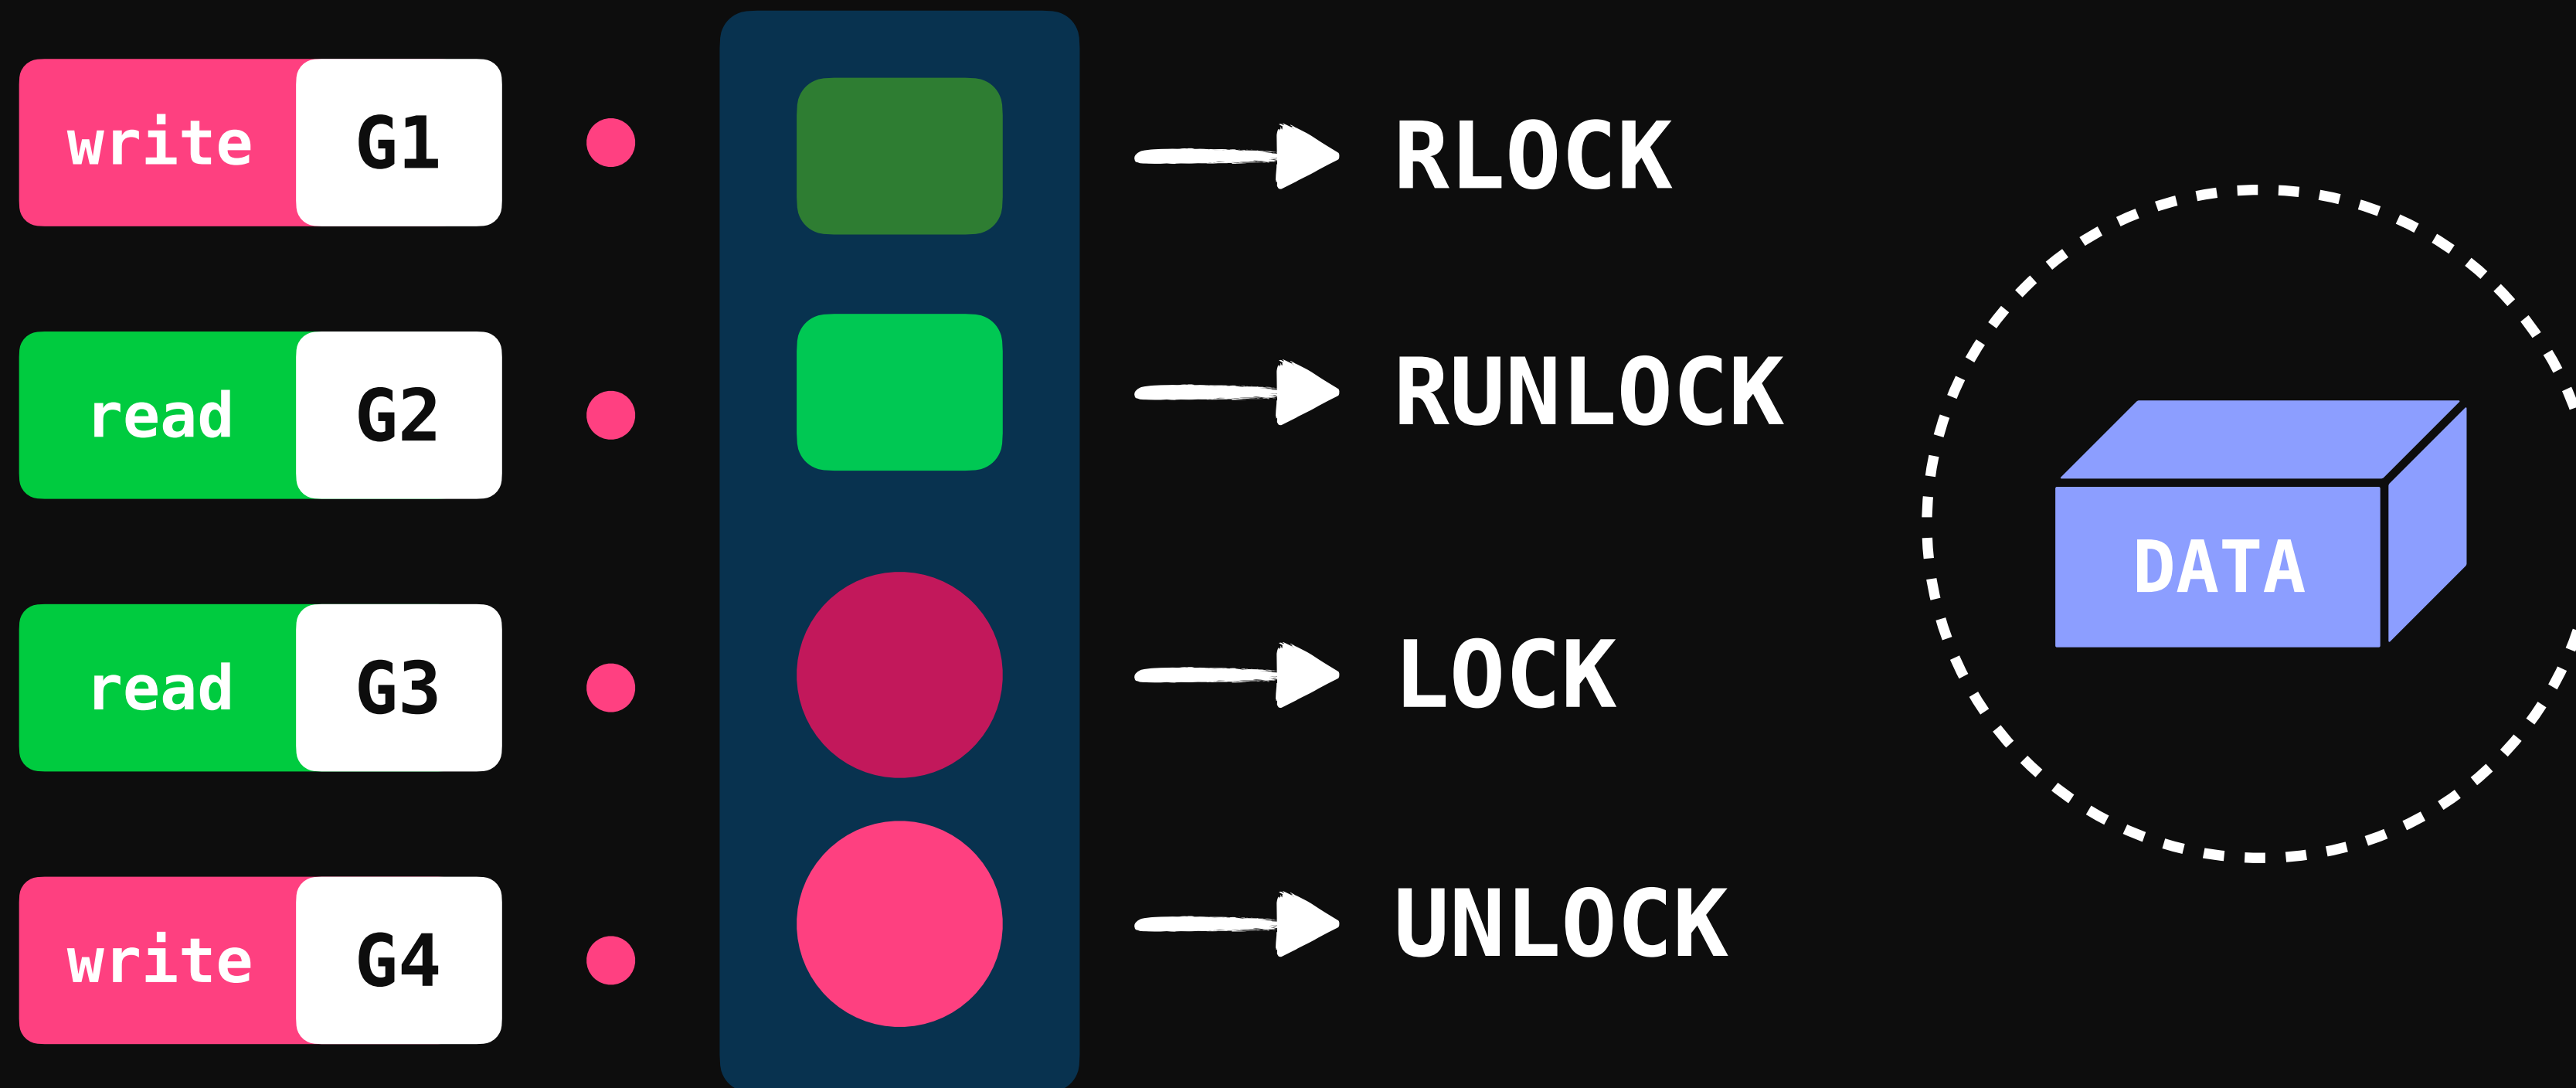
MUTEX

WRITE ONCE, READ MANY

RWMUTEX



RWMutex



G1

G2

G3

i++

i++

i++

var i

i could be 1

i could be 2

i could be 3



`i++`

`get value of i`

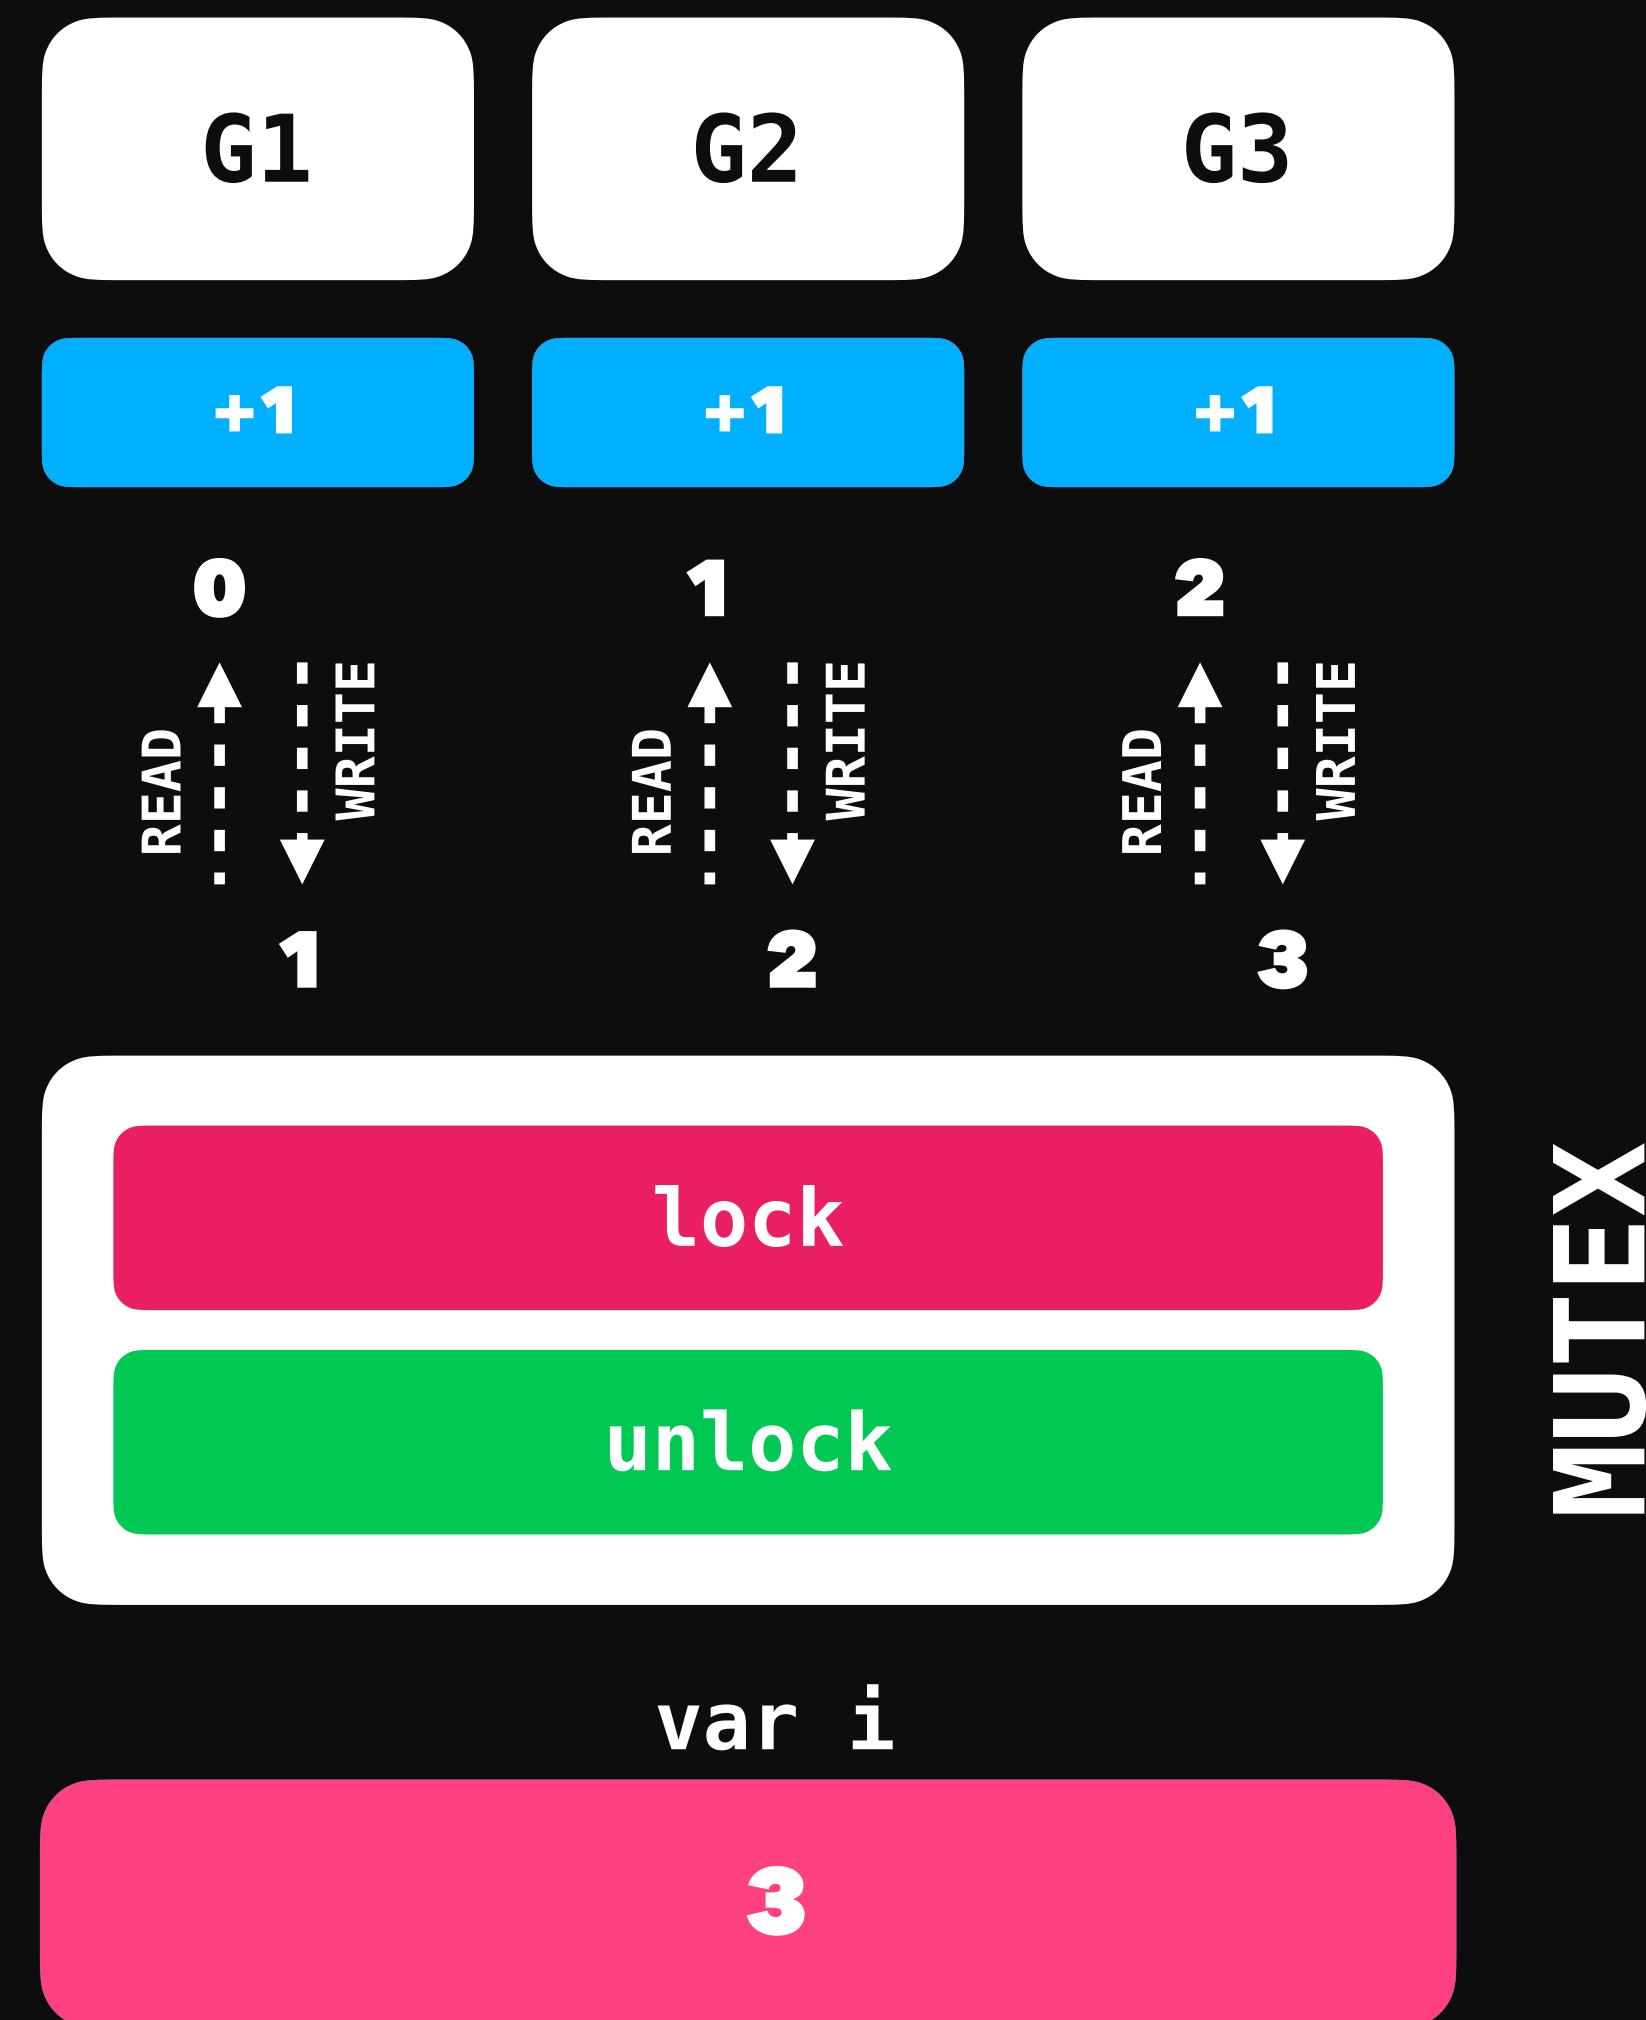
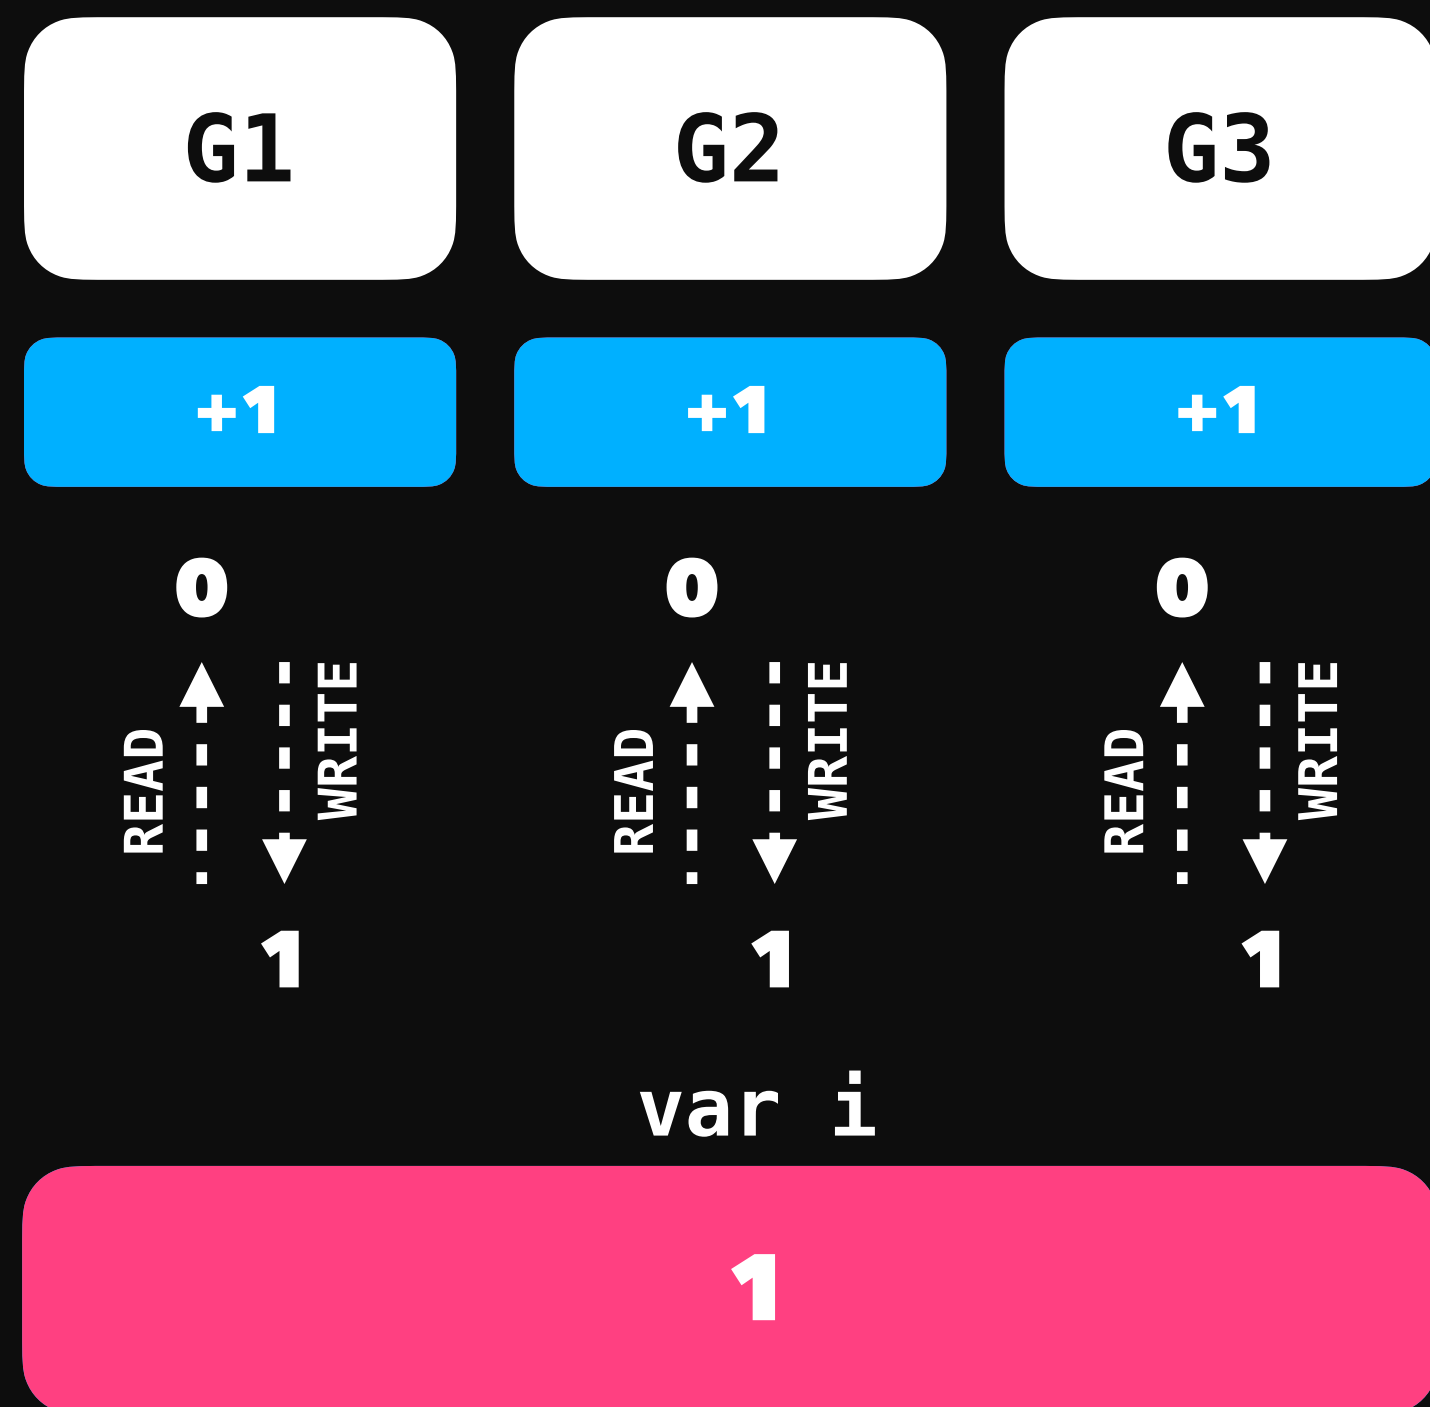
`increment value of i`

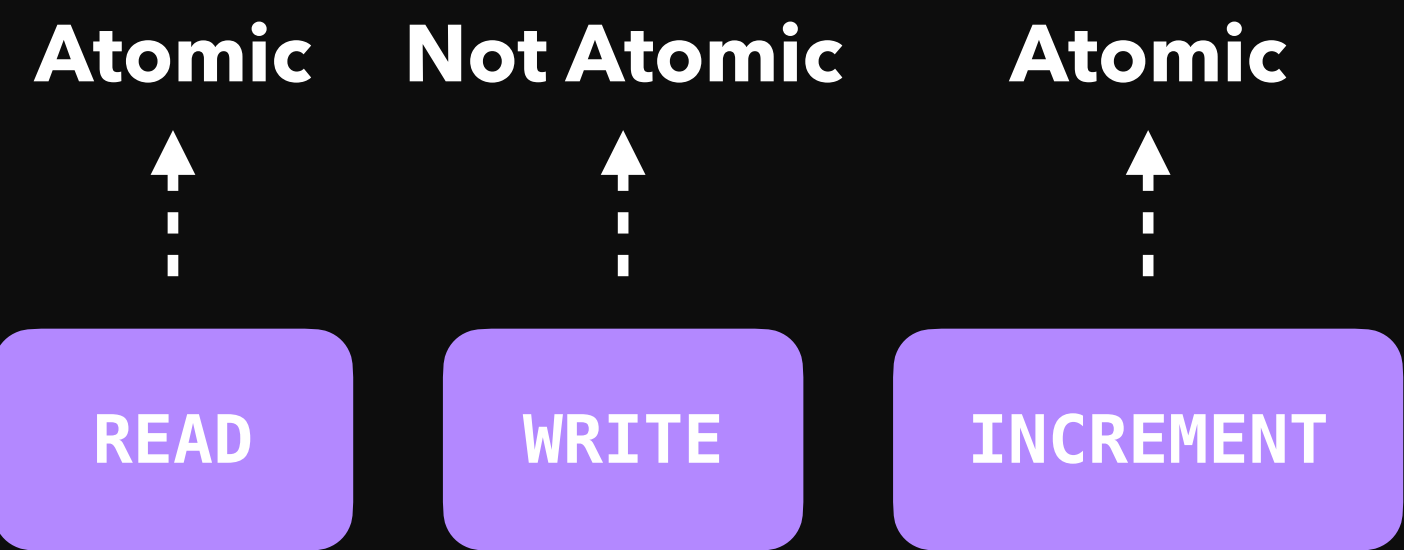
`store value of i`

INDIVISIBLE

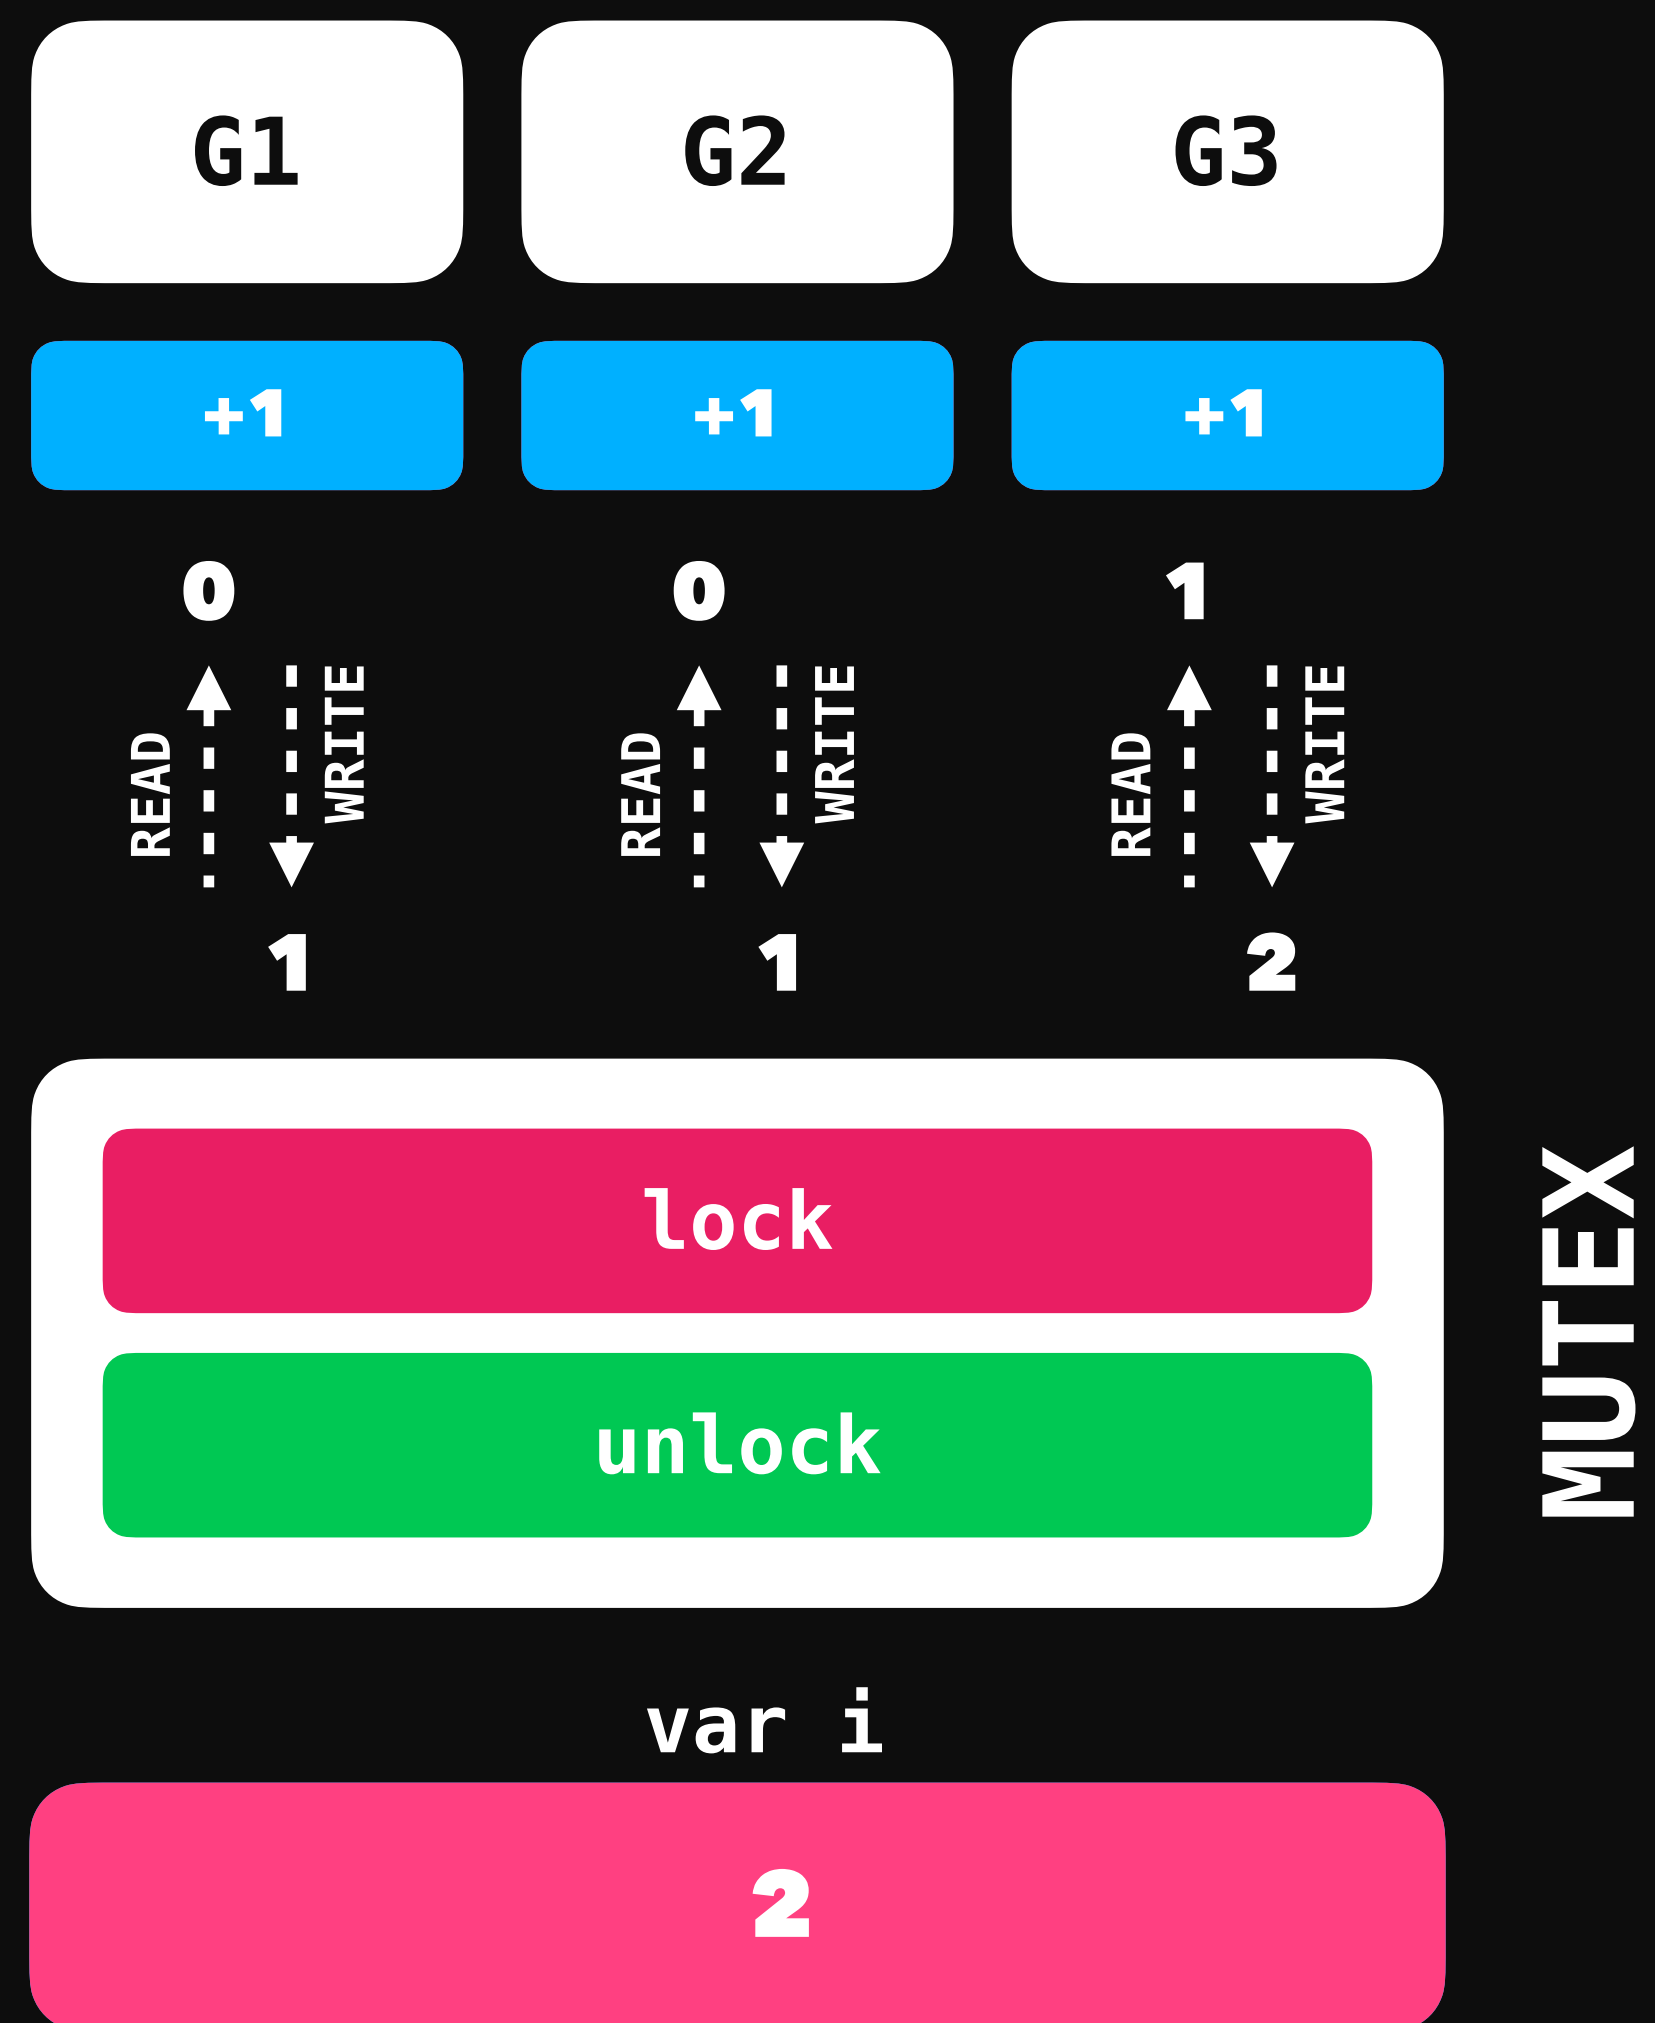
UNINTERRUPTIBLE

ATOMIC





IS THAT SO?




Single Go Routine Context




NO MUTEX

Main Go Routine Context



NEEDS MUTEX


Multiple Go Routines Read Only Context



NO MUTEX

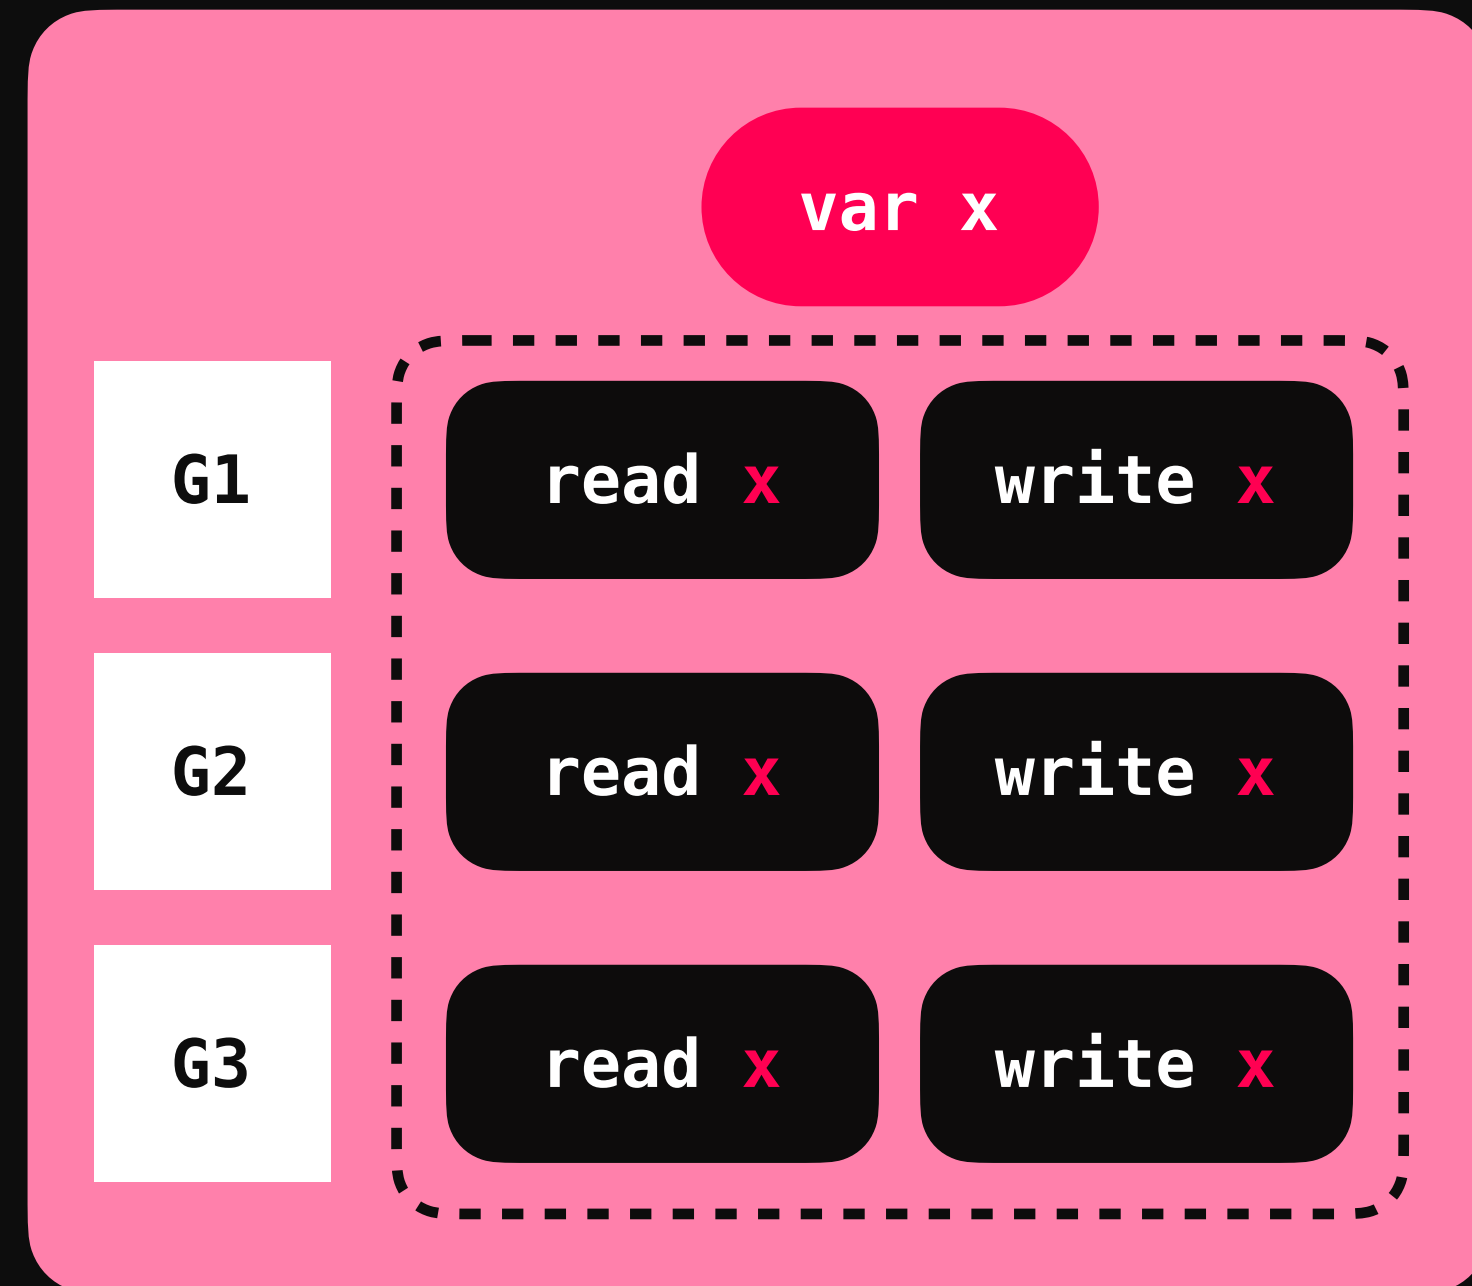
Multiple Go Routines

Read Write Context

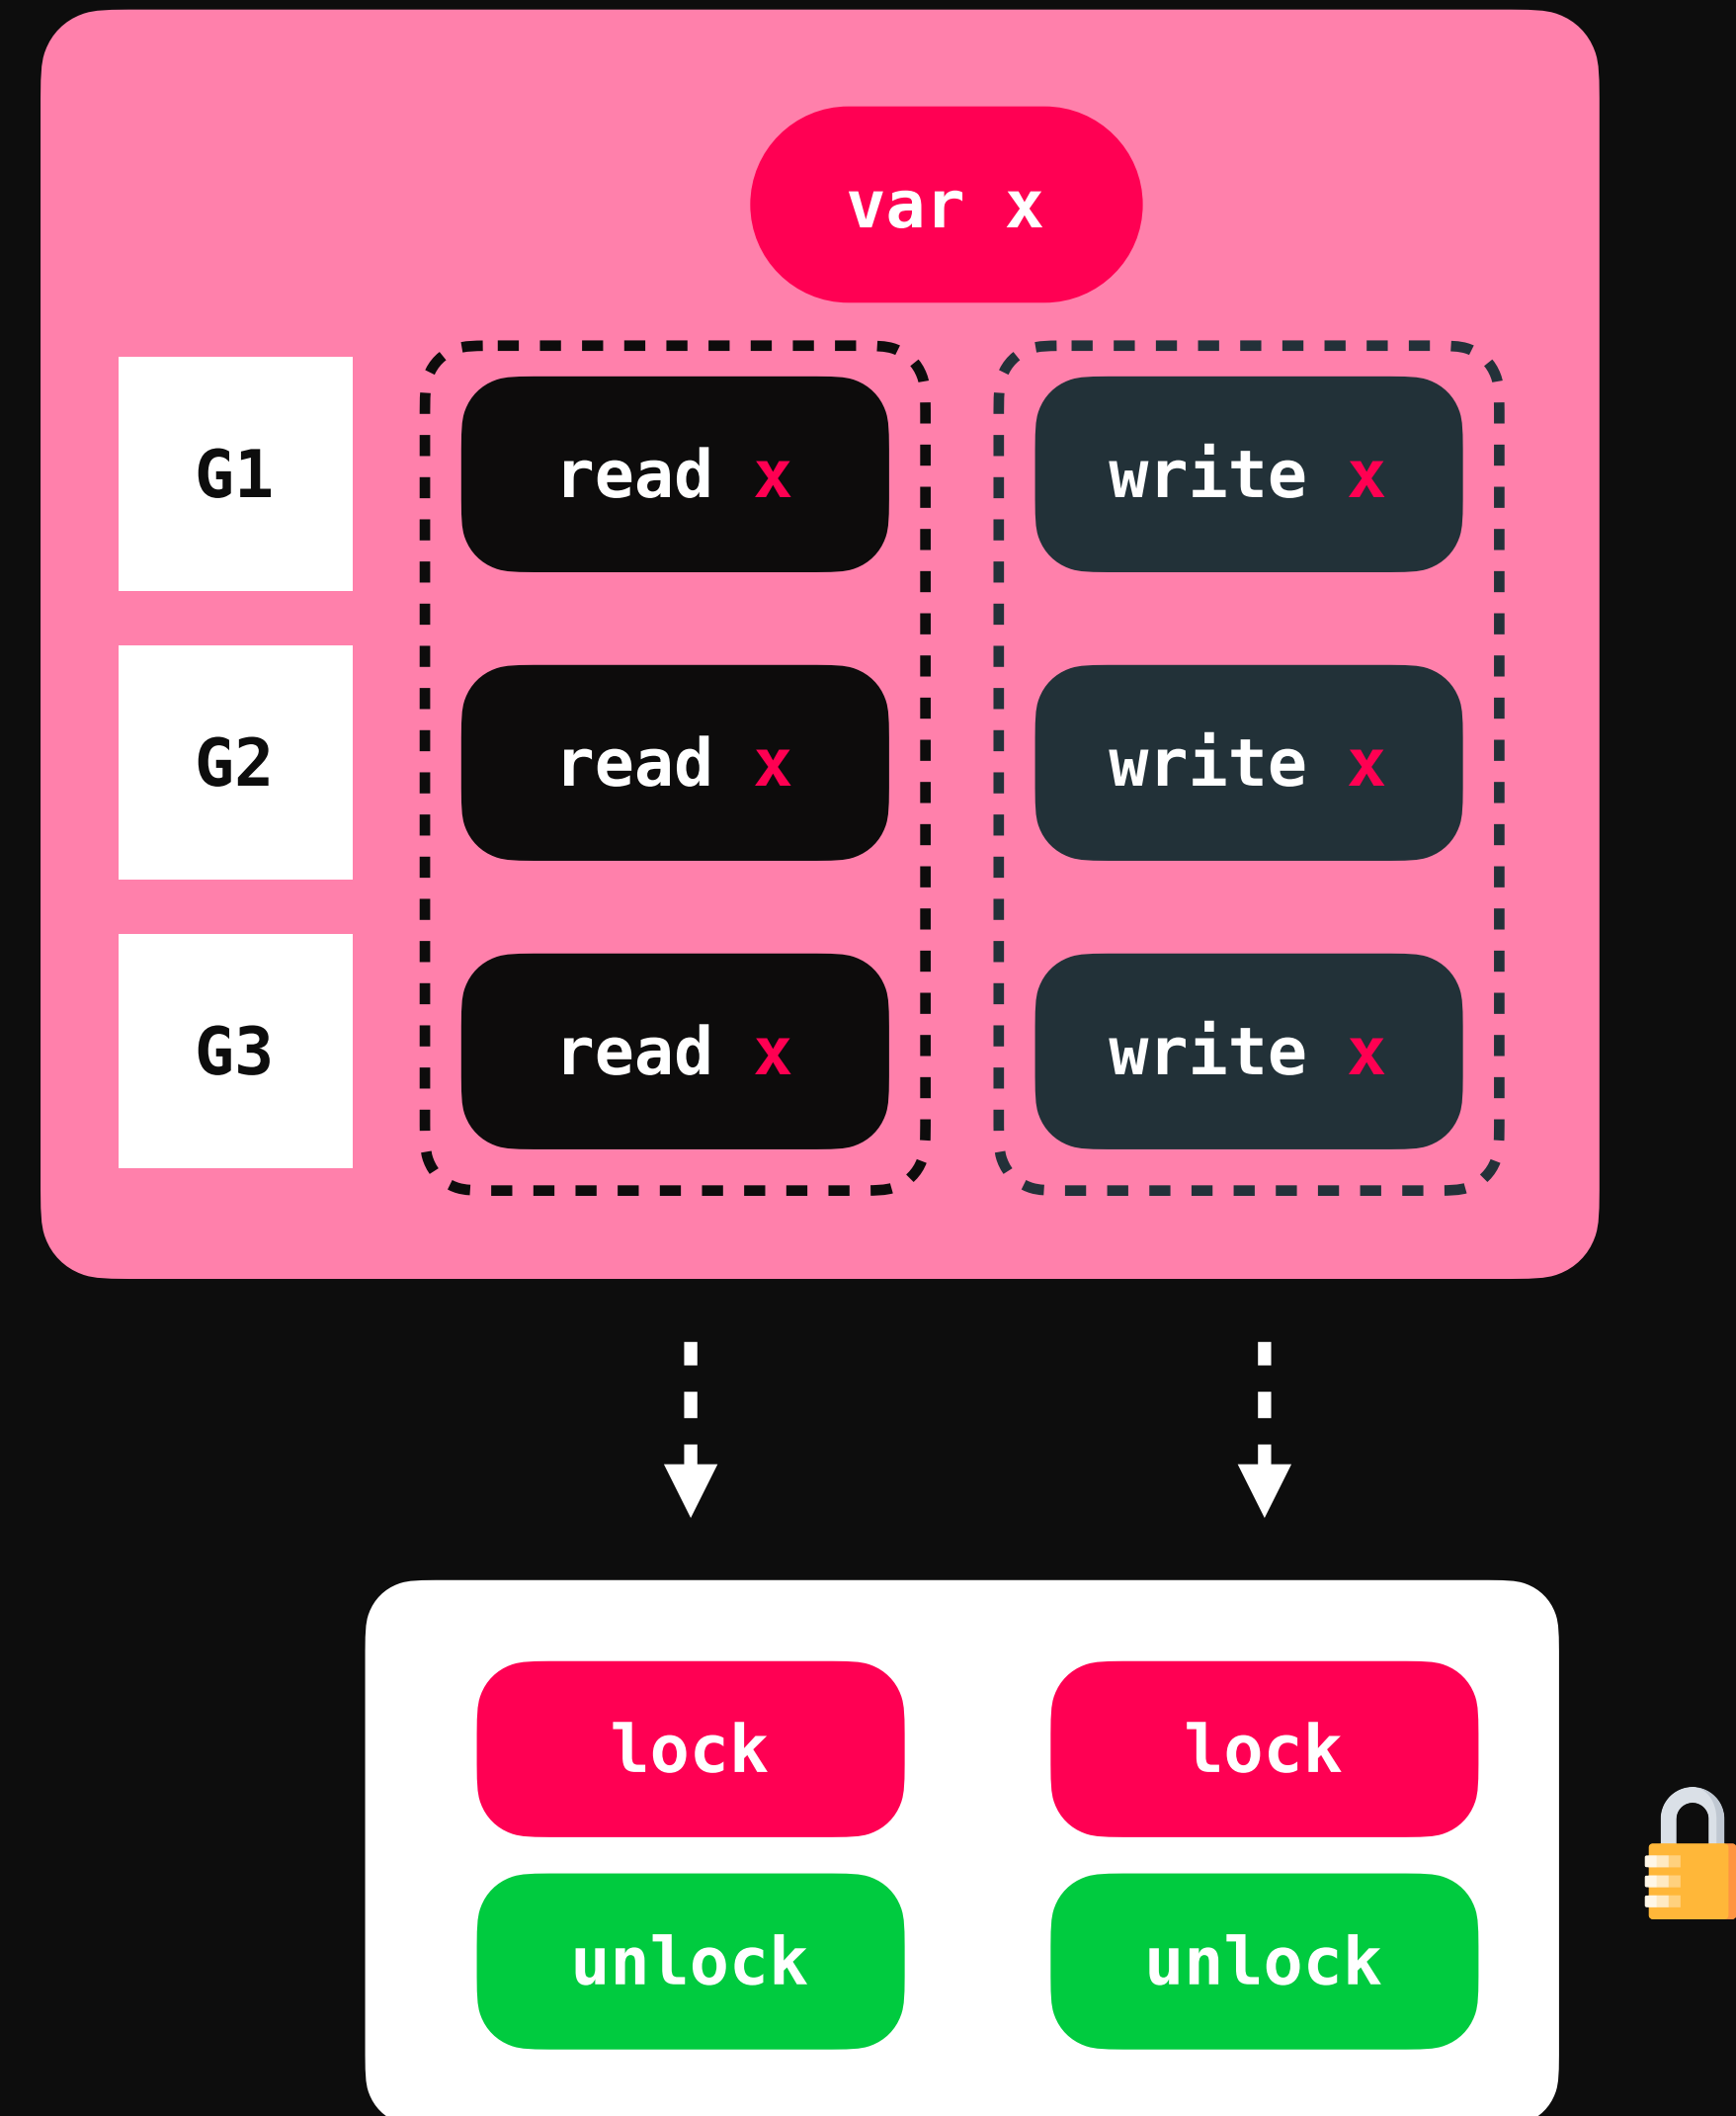


NEEDS MUTEX

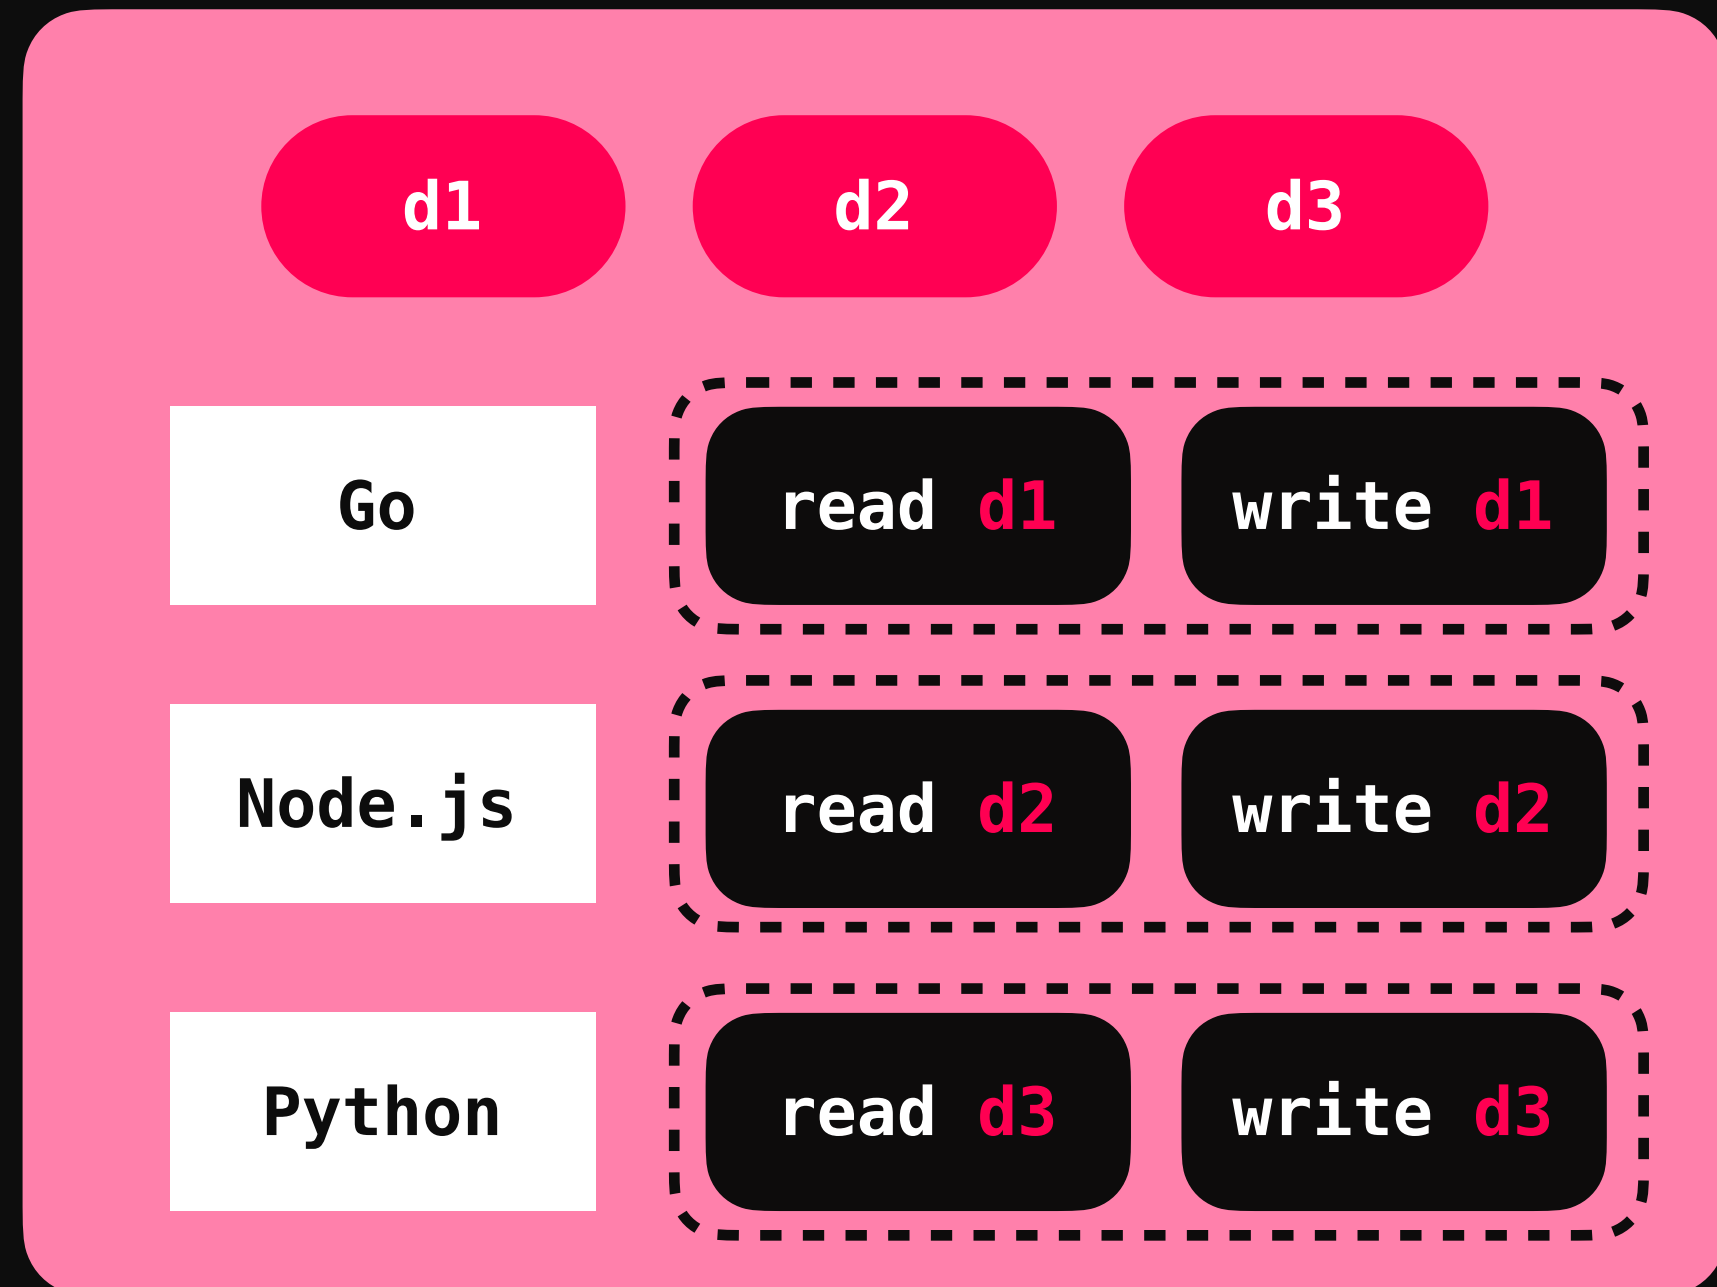
Coarse Grained Context



Fine Grained Context




OS Context Different Locations





NO LOCK



OS Context Same Location



NEEDS LOCK





MUTEX LOCK

⋮

⋮

G1

lock()

unlock()

G2

lock()

unlock()



DEADLOCK

COFFMAN CONDITIONS



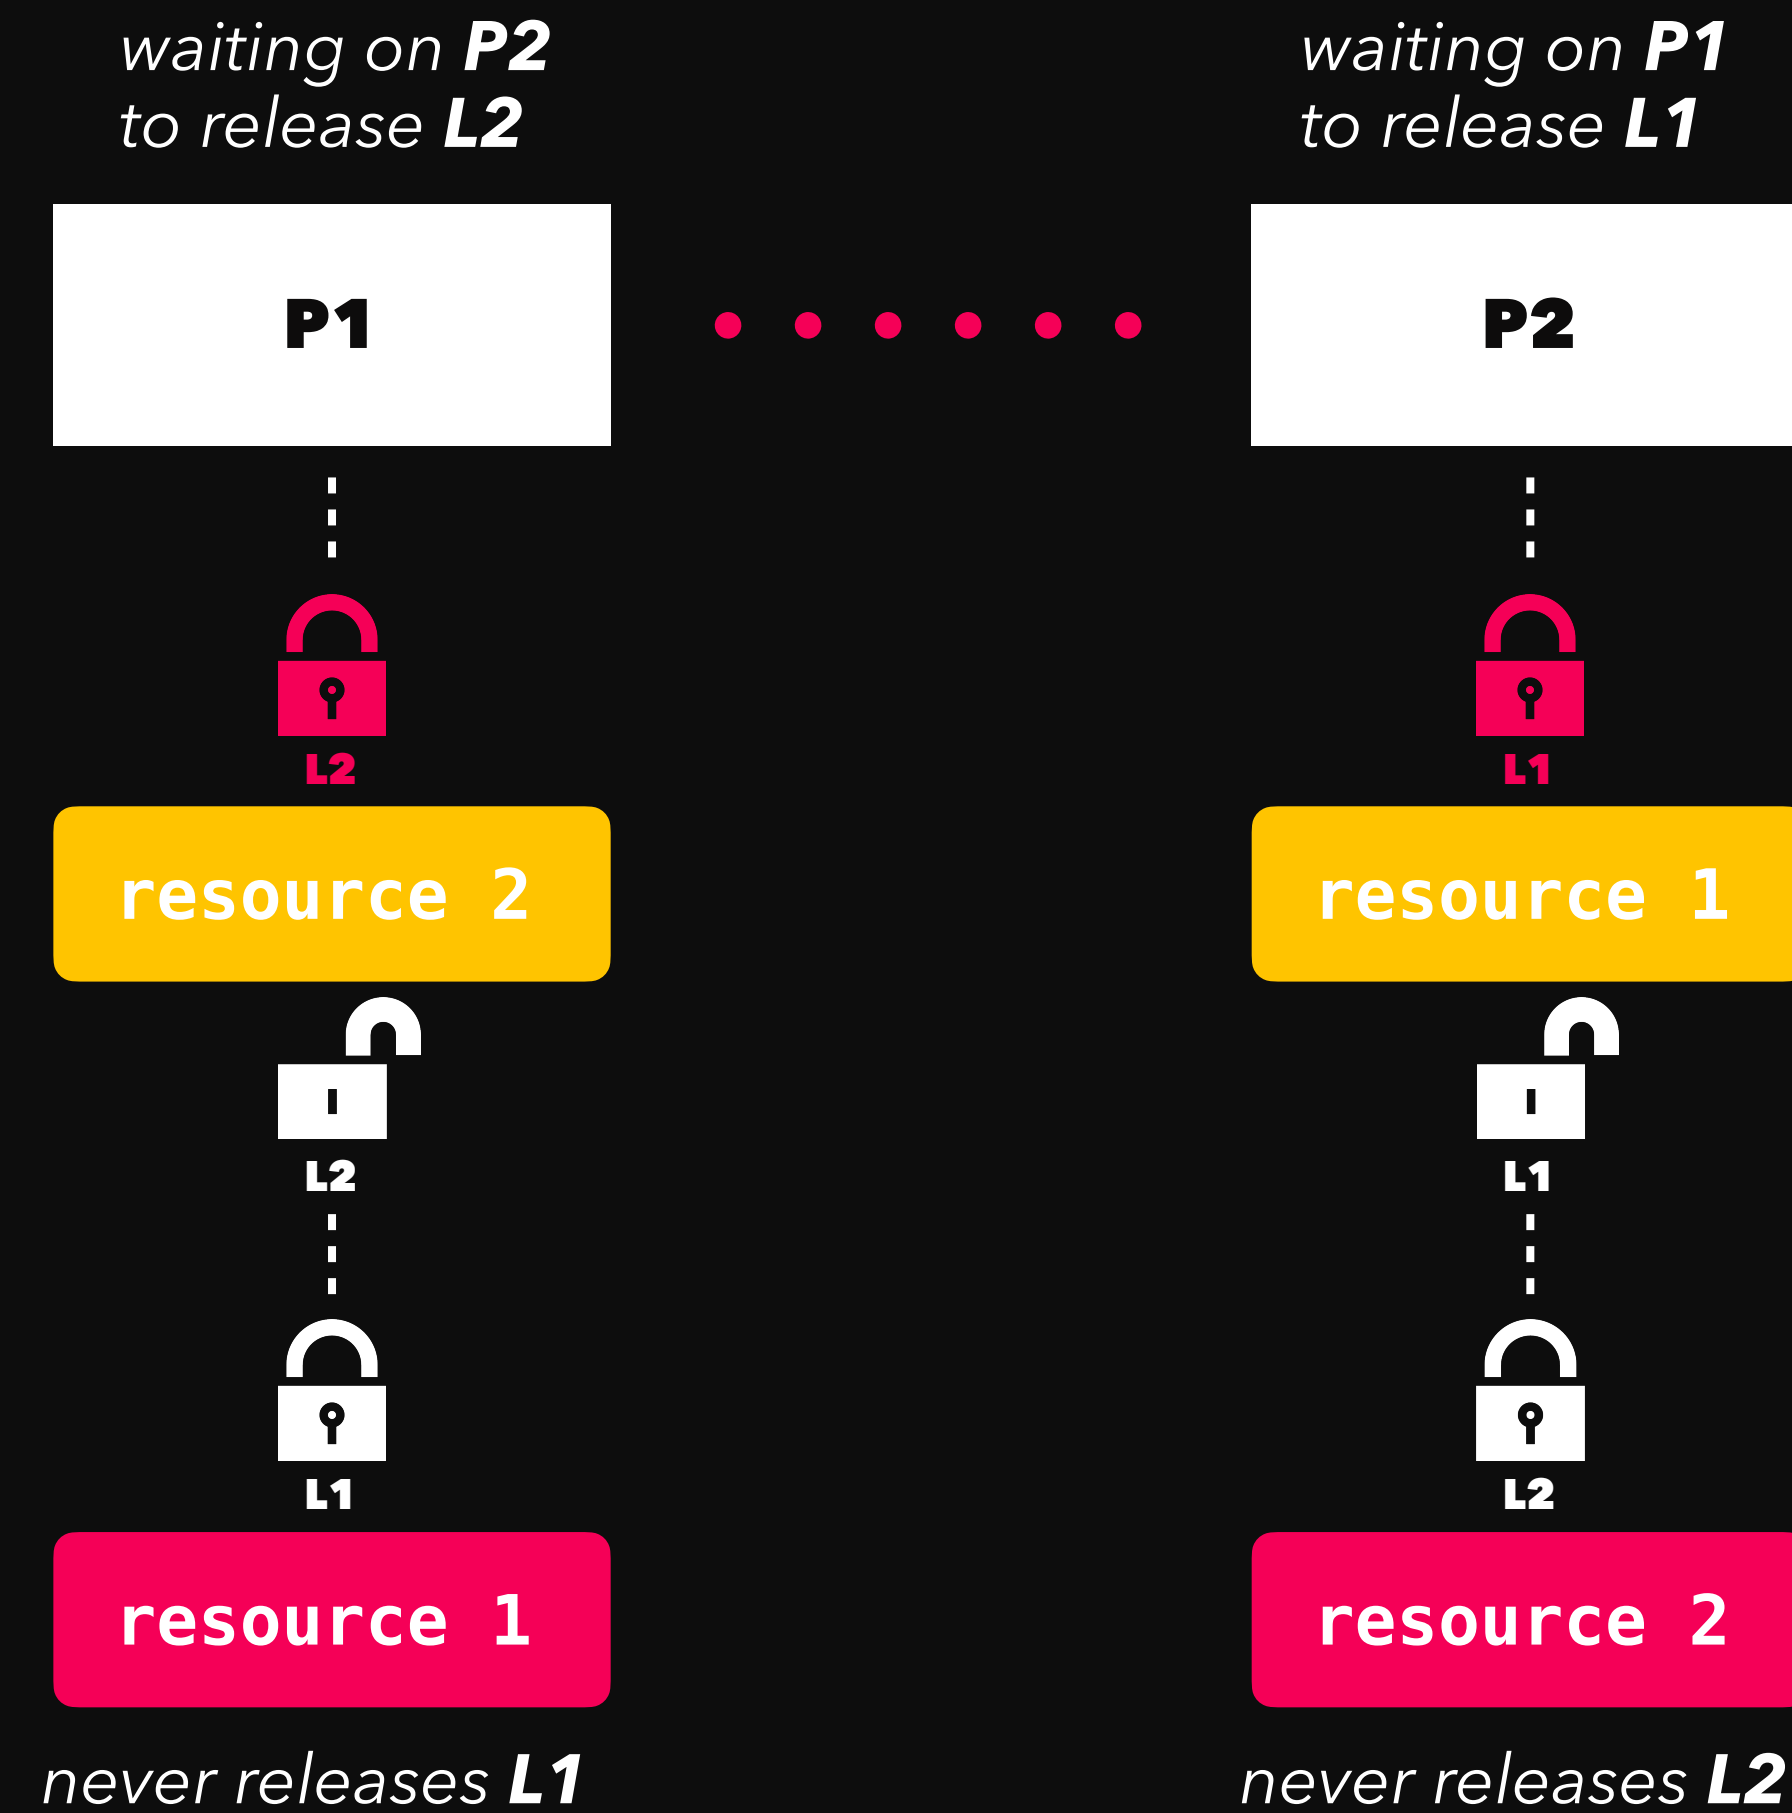
WHAT IS A DEADLOCK?



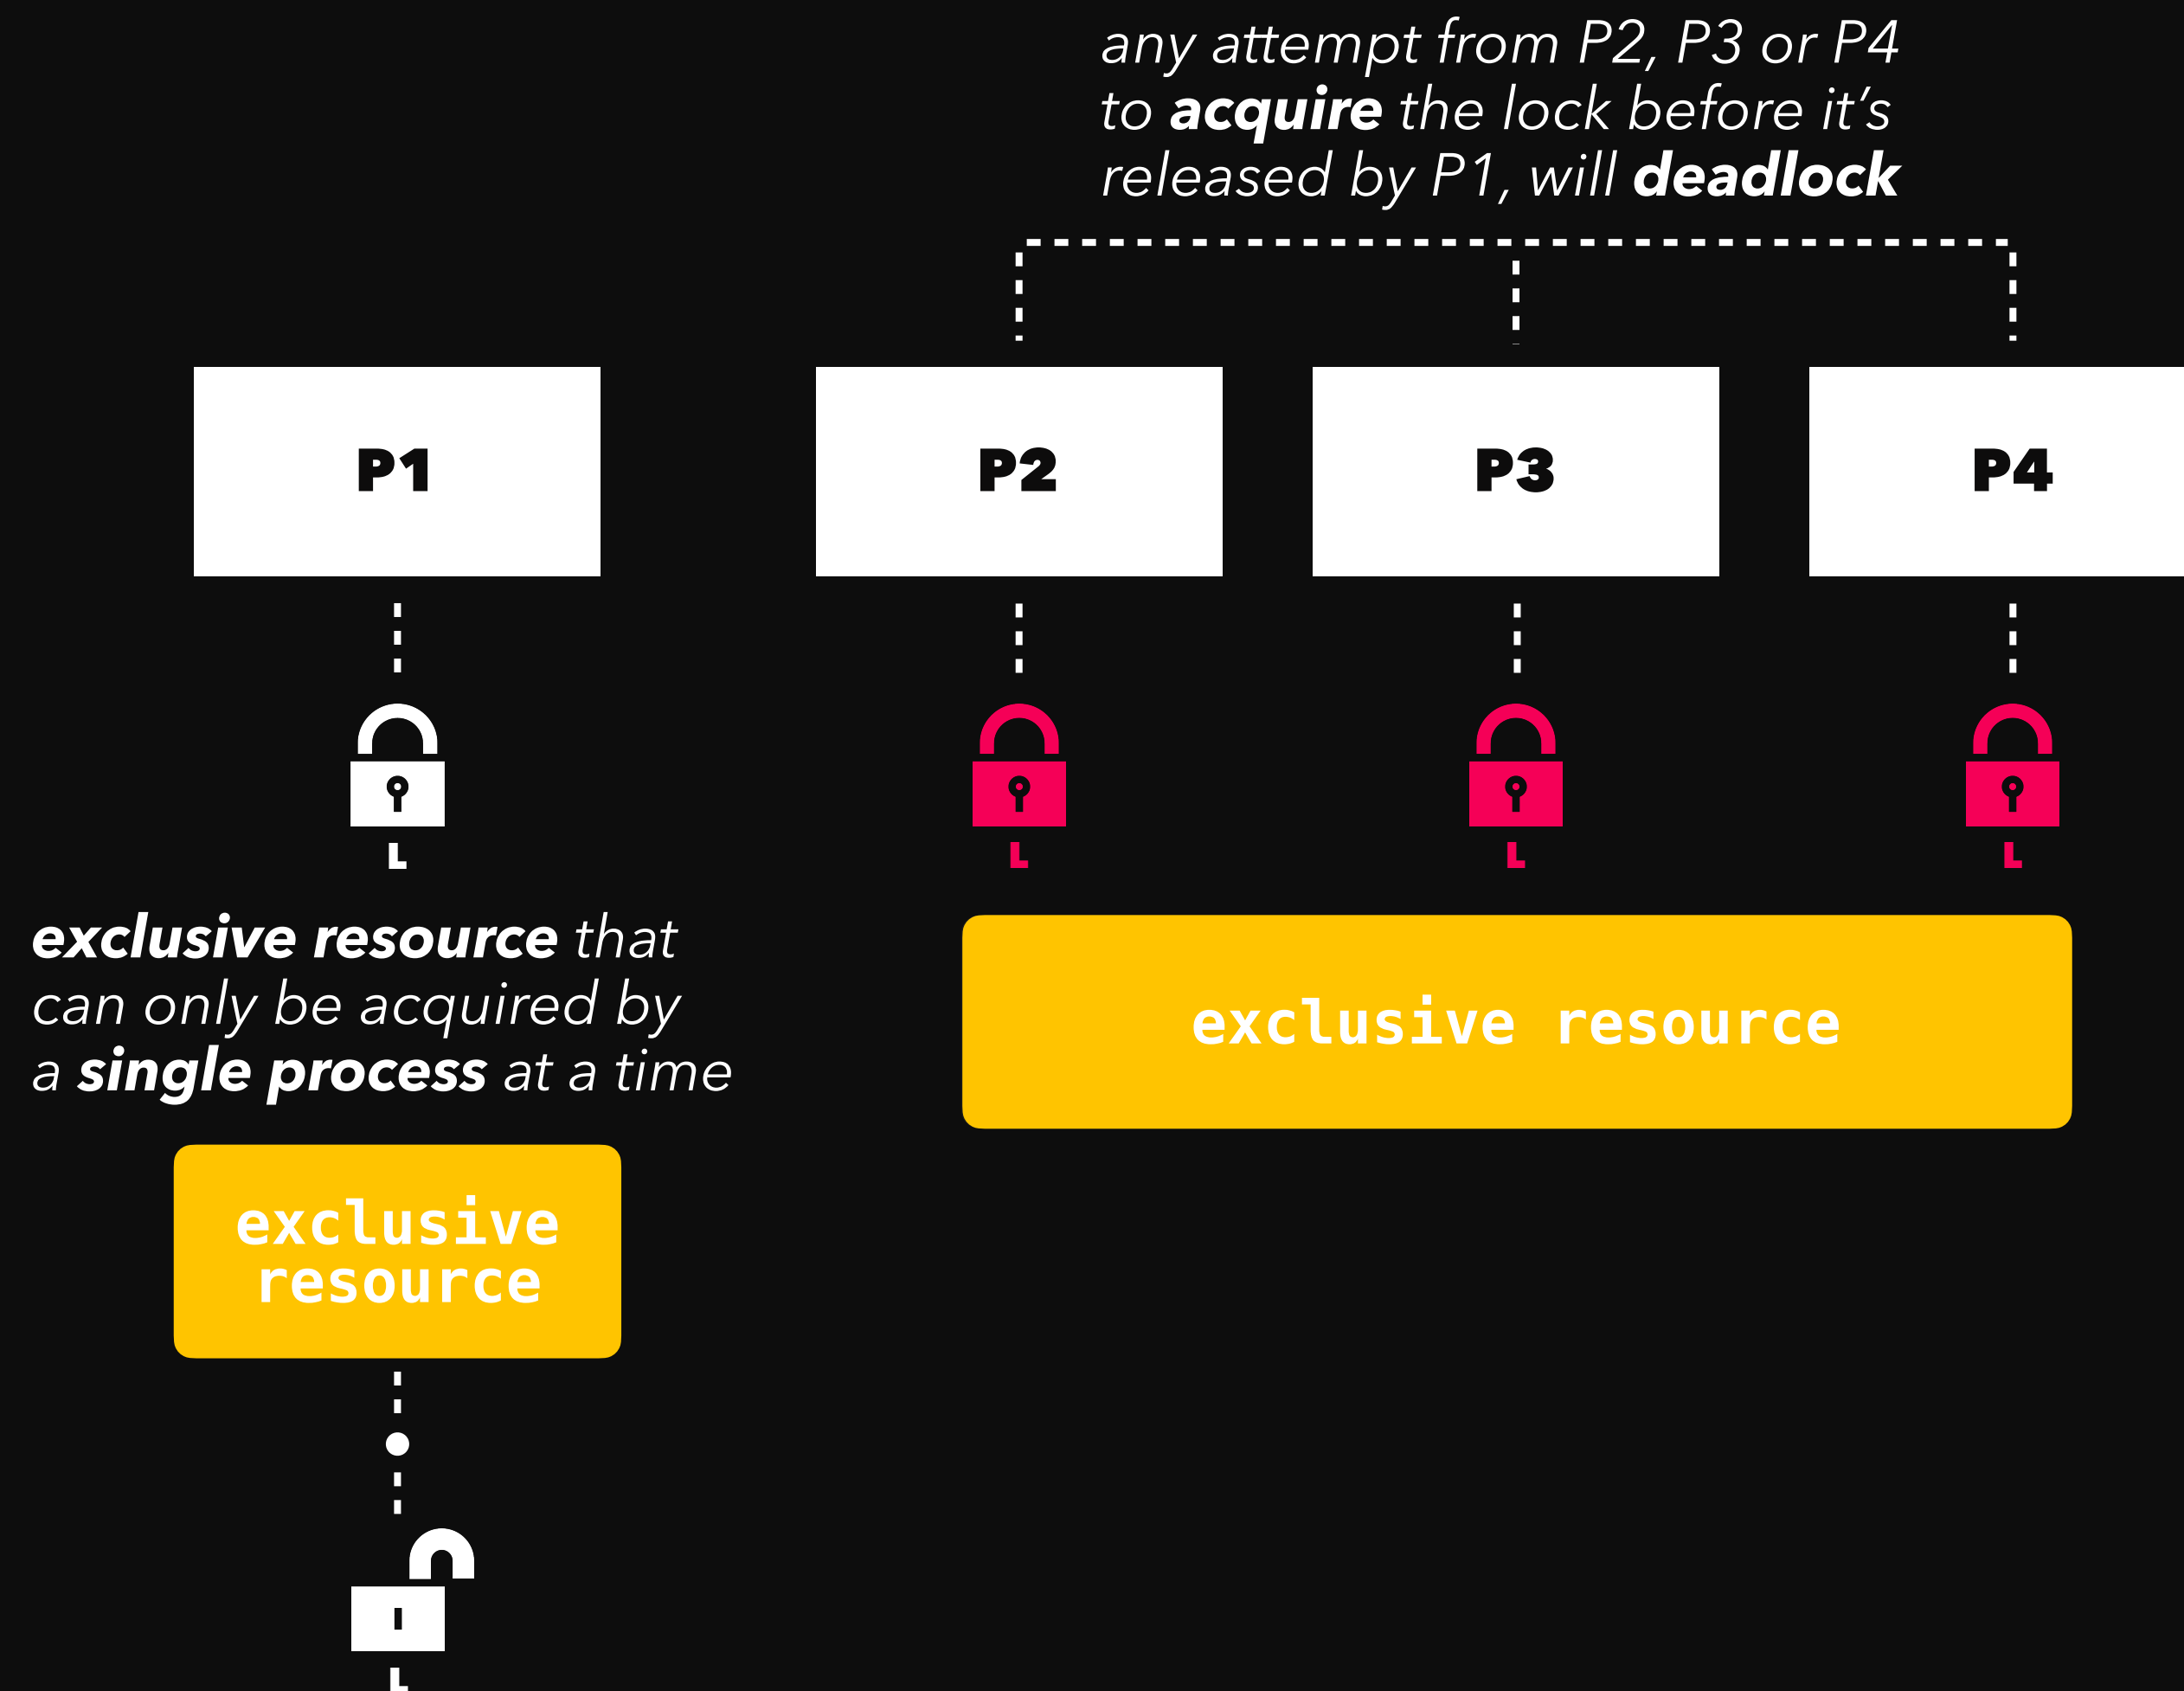
*A **deadlock** is a state in which each member of the group **waits** for another member, including itself, to take action, such as sending a message or more commonly **releasing a lock**.*



1 CIRCULAR WAIT




2 MUTUAL EXCLUSION



3 HOLD AND WAIT

condition inside P1 is never satisfied



waiting on P1 to satisfy the condition



WHAT IS PREEMPTION?

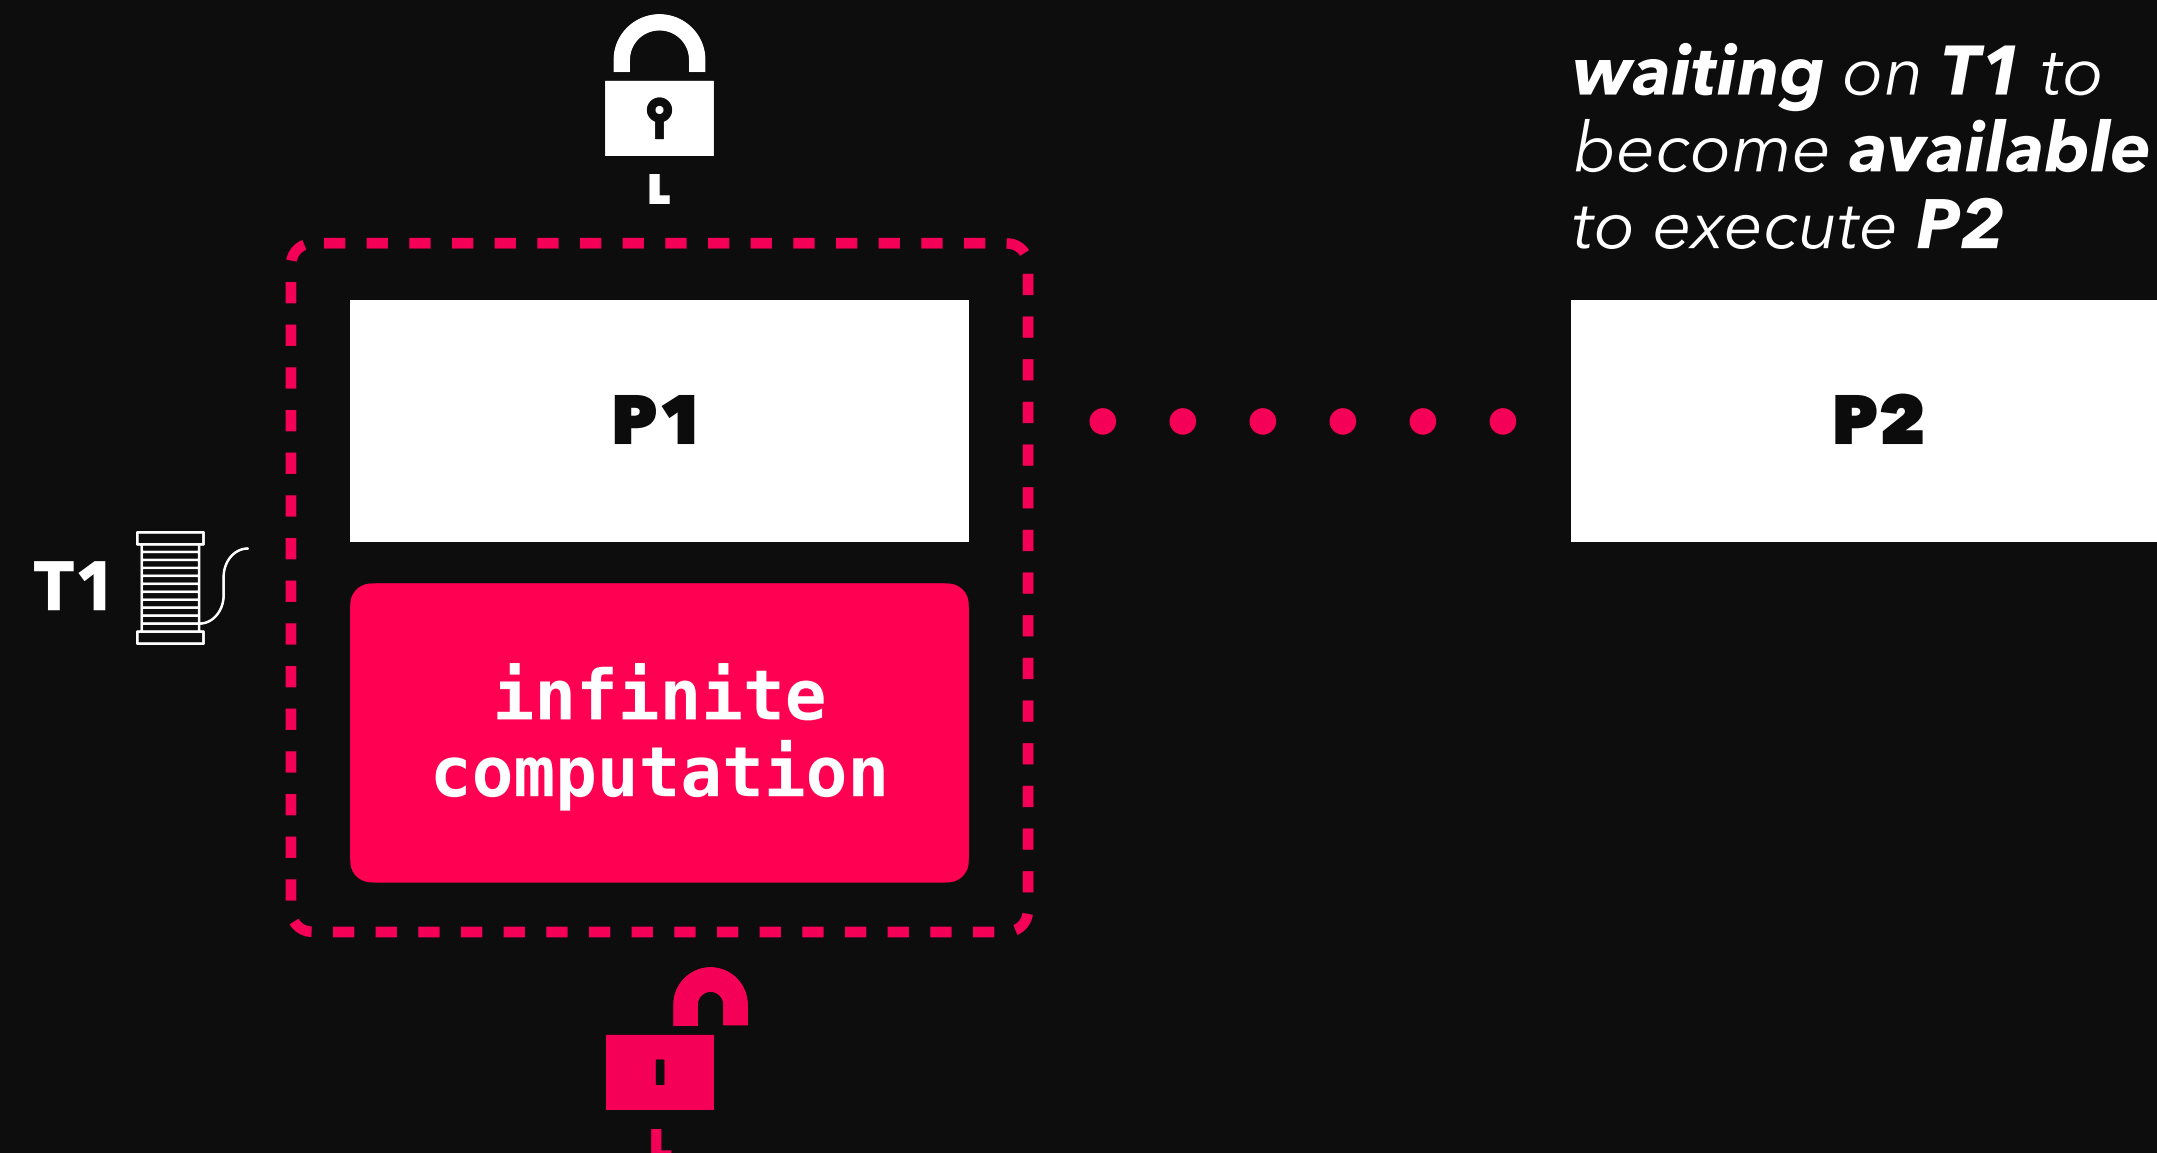


***Preemption** is the act of temporarily **interrupting** an **executing task**, with the intention of **resuming** it at a **later time**.*

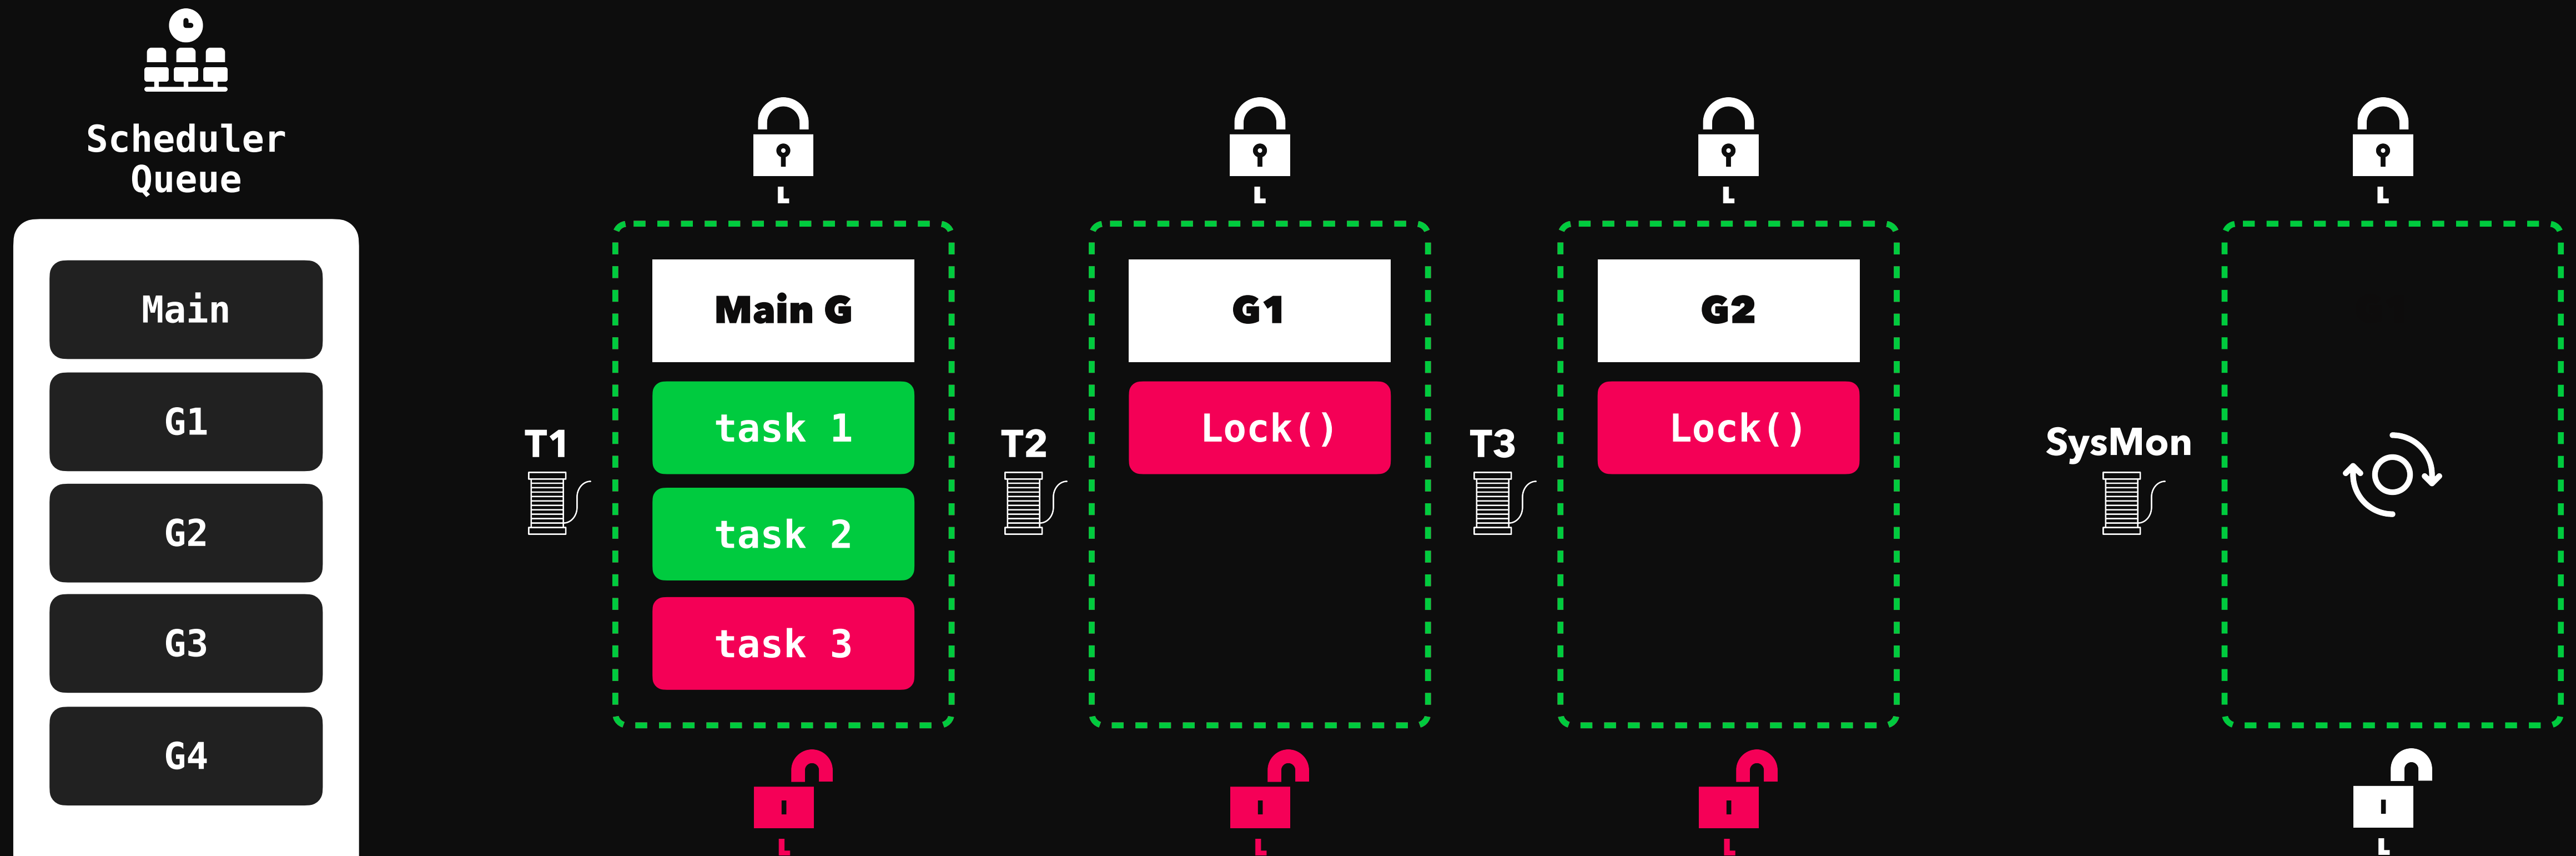


4 NO PREEMPTION

***T1** is always **stuck** executing **P1**, thus it's **never available** for other processes*



CHECK DEAD



all go routines are asleep – deadlock

Mike



17m

Jane





30m


John



15m



Mutex



released

G1

released

G2

G1 work

G2 work





3



REFILL



1



MUTEX LOCK

⋮

⋮

G1

acquire

work

sleep

release

100μs

G2

sleep

acquire

work

release

100μs



Scheduler Queue

G2



10ms

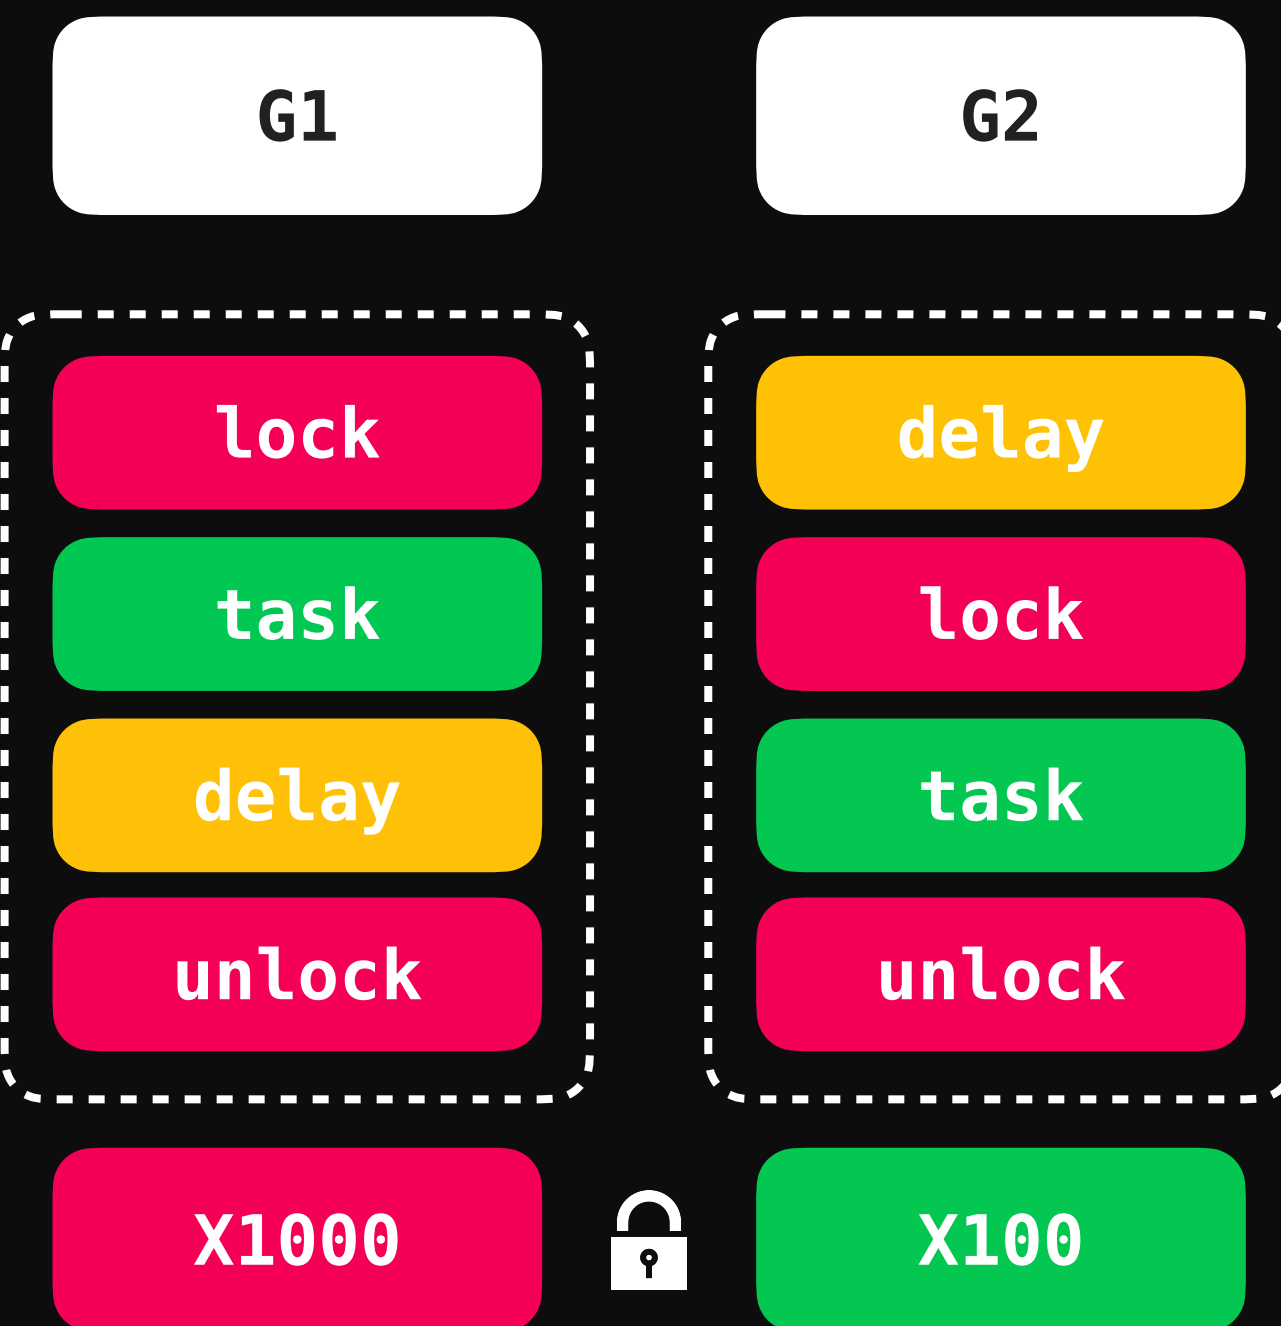
CONTENTION




● UNEVEN WORK

VS

STARVATION



● UNEVEN TIMING/HOLD TIME





MUTEX LOCK

G1


lock()

condition

unlock()

wait()

John



var left

var right

Hallway

G2

lock()

condition

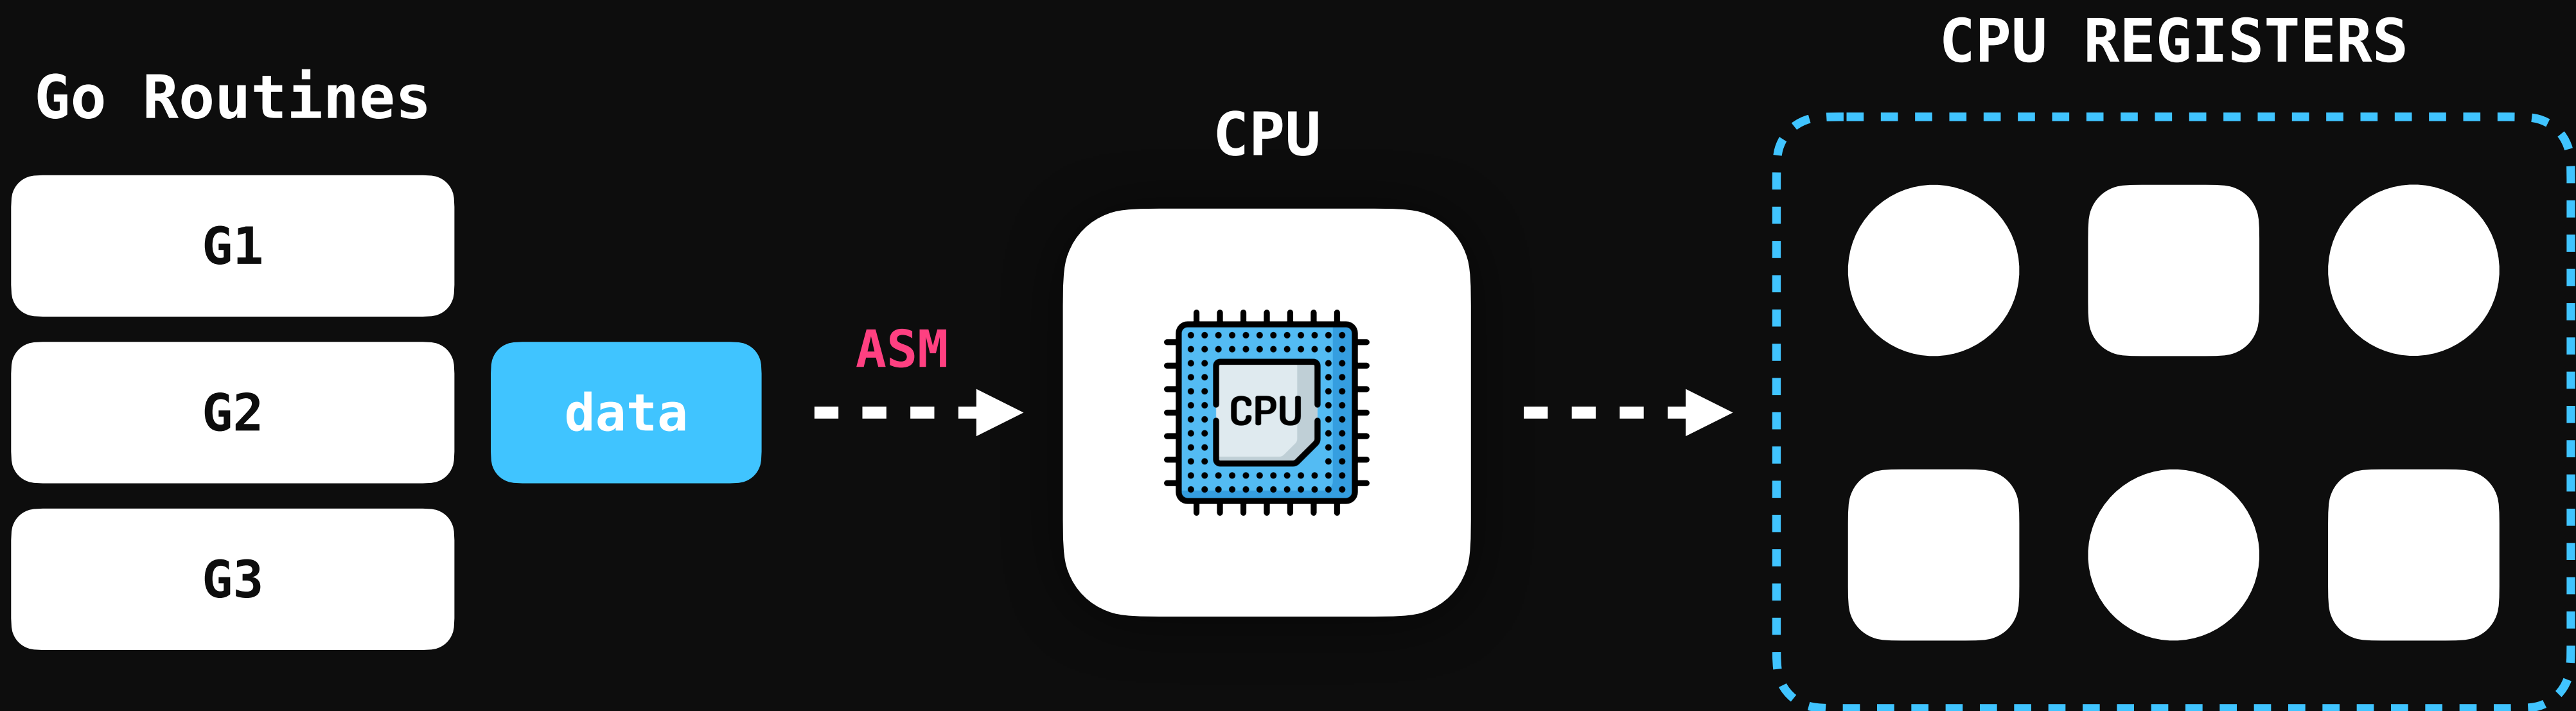
unlock()


wait()

Alice

sync.Locker

```
type Locker interface {  
    Lock()  
    Unlock()  
}
```





Lock()
Unlock()

TEST & SET

Prefer **Atomics** over **Mutexes** for **simple data**

Use **Mutex** for write heavy scenarios

Use **RWMutex** for read heavy / mixed scenarios

Prefer **Fine Grained Context** where possible (limit the context)

Don't use the mutex **longer** than you need to (**avoid extra hold time**)


Use **types**, a **local mutex** and **methods** over direct mutex calls

Avoid **Contention** by **distributing work evenly**

Avoid **Starvation** by testing for **mutex fairness**

Use **sync.Locker** for generic code that uses a mutex

Concurrency Control ensures that **correct results** for **concurrent operations** are generated, while getting those results as **quickly** as possible.



ACID

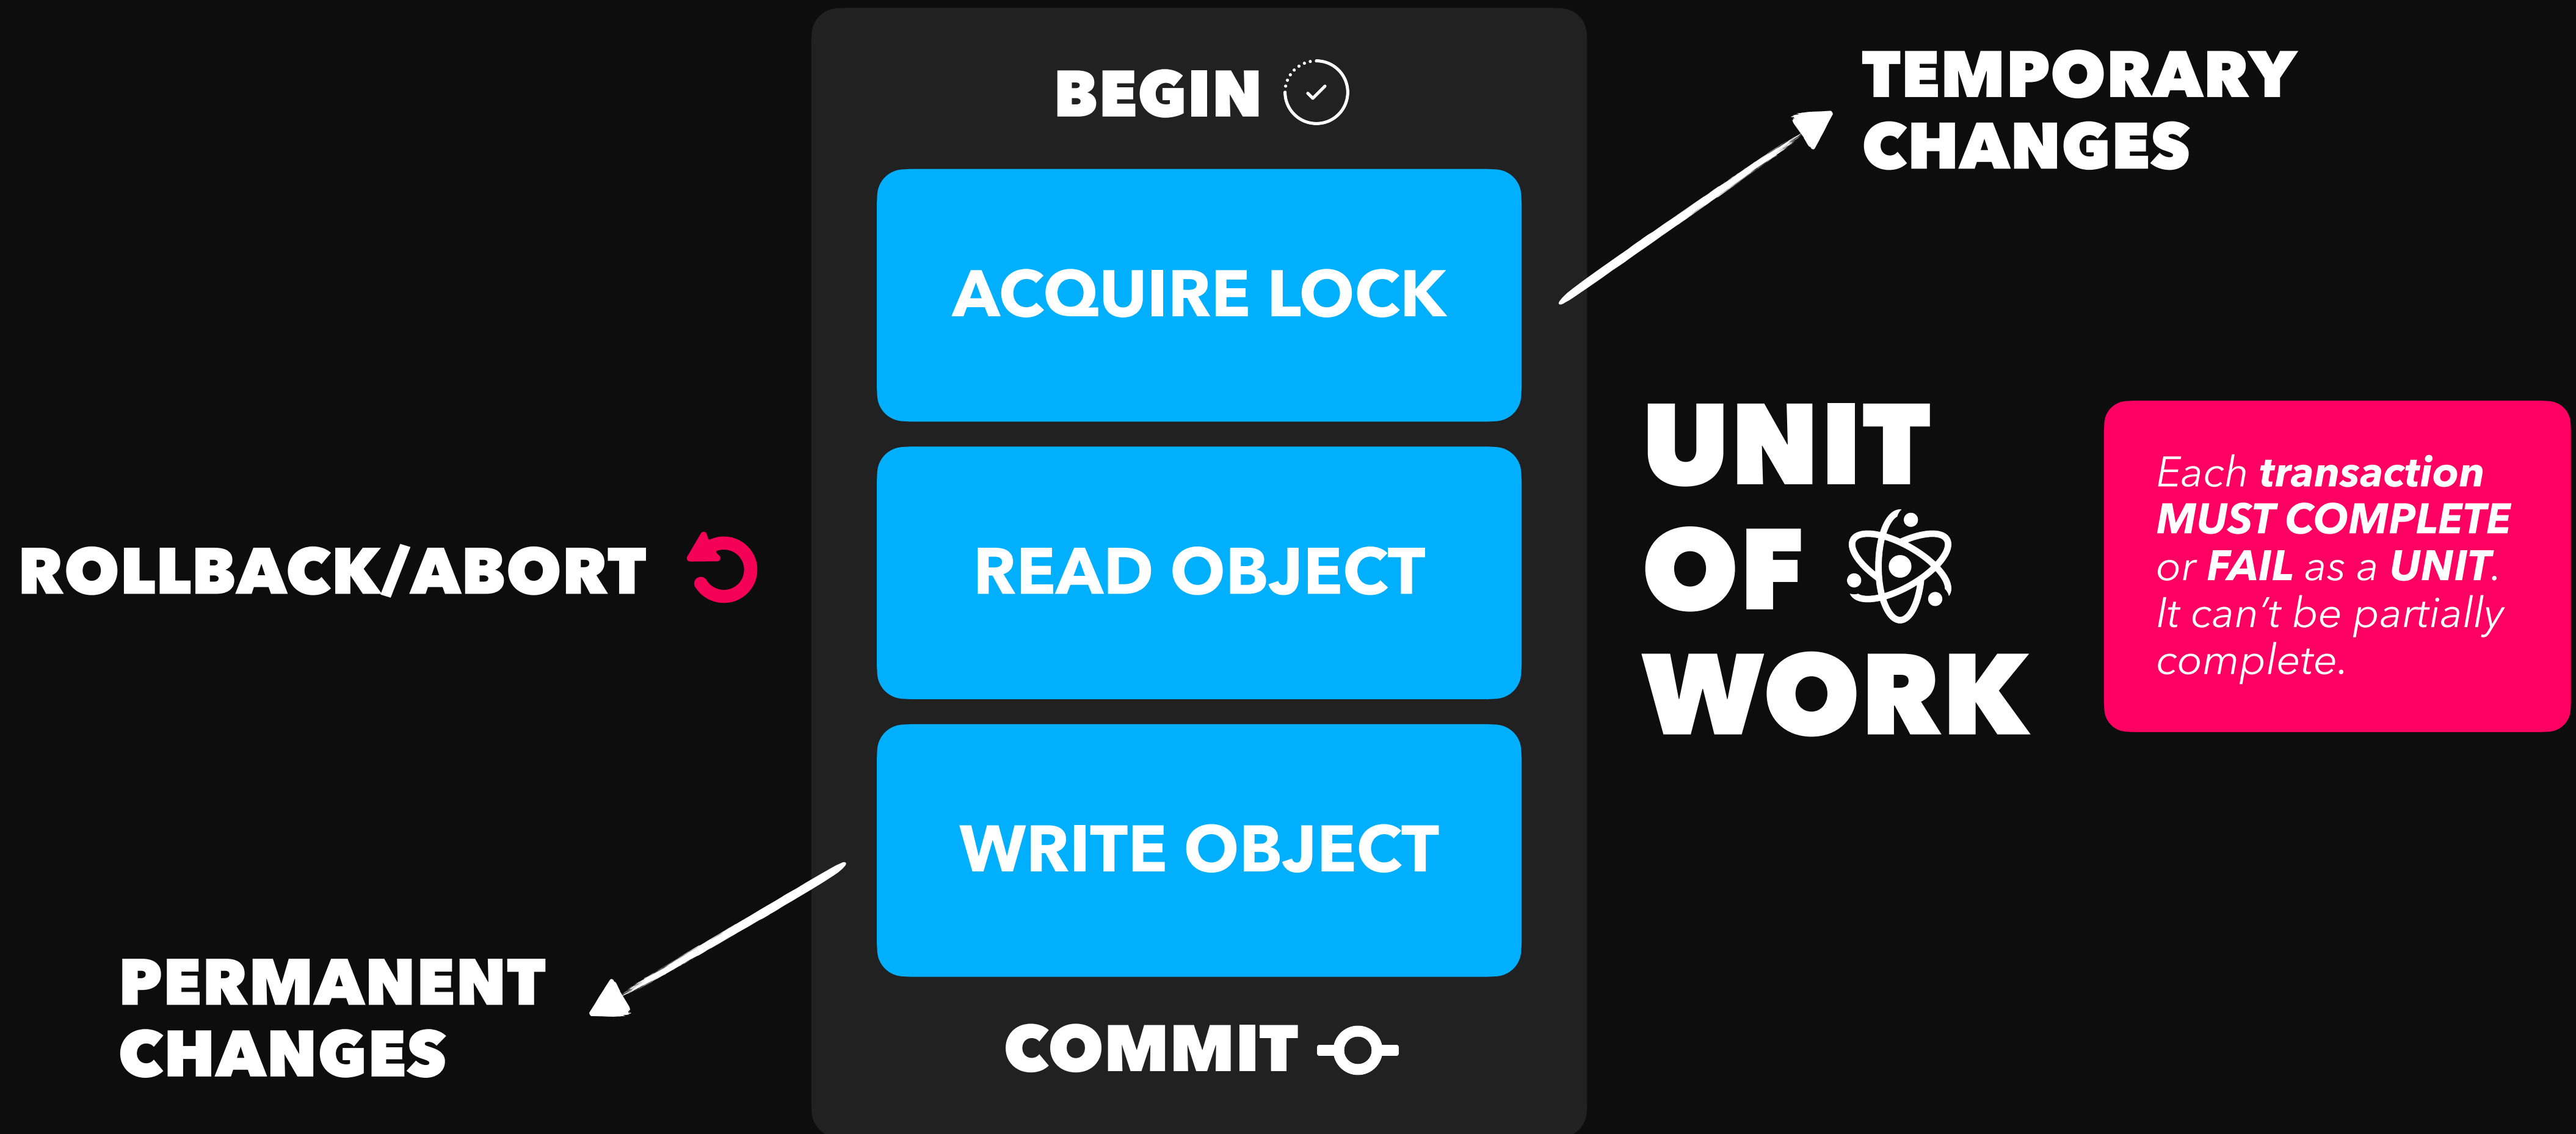
ATOMICITY

CONSISTENCY

ISOLATION

DURABILITY

TRANSACTION





\$60
-\$40



Alice




\$40



Book Store

No **payment**
without the **book**

No **book** without
successful payment



Alice wants to initiate
2 **transactions** at the
same time, transferring
John **\$100/transaction**



\$200



\$0
-\$150



\$200



Alice

TX1

\$100



TX2

\$100



John

Alice wants to initiate
2 **transactions** at the
same time, transferring
John **\$100/transaction**



Alice wants to initiate
2 **transactions** at the
same time, transferring
John **\$100/transaction**



\$150



\$50

-\$100



\$100



Alice

TX1 → \$100


TX2 → \$100




John

1d

In **Concurrency Control** of databases and various **transactional applications**, a transaction schedule is **Serializable** if its **outcome** is **equal** to the outcome of its transactions **executed serially**, without overlapping in time.




SCHEDULING



SERIAL

NON-SERIAL



T1

T2

Serial Schedule

T1

T2

Non-Serial Schedule

R

W

R

W

R

W

W

R

T1

R

W

R

W

T2

R

W

W

R

R1

W1

R2

W2

Operations are **executed**
in an **interleaved** manner

R1

R1

R2

R2

R2

W1

R1

R1

W1

R2

W2

W1

W2

W2

W1

W2

• S1

• S2

• S3


• S4

CORRECT, BUT SLOW

CORRECT & FAST
SERIALIZABLE


FAST, BUT INCORRECT

SERIALIZABILITY



RESULT EQUIVALENT

CONFLICT SERIALIZABLE




VIEW SERIALIZABLE

RESULT EQUIVALENT



CONFLICT SERIALIZABLE




Conflict operations are on the **same data item** i.e **A, B**

RW, WR, WW conflicts MUST be on **different** transactions


At least **1** of conflict operations MUST be a **Write**

2 READ operations will **not** create a **conflict**




VIEW SERIALIZABLE

S1



SERIAL 

S2



NON-SERIAL 

 *FIRST **READ** performed by the **same transaction***

Initial operations on each **data item** are on the same transaction

Final operations on each **data item** are on the same transaction


 *LAST **WRITE** performed by the **same transaction***

Number of **WR** operations MUST be the **same** on all schedules

 *Maintain **Producer-Consumer/W→R** sequence*

S1

- **Initial operation** of data item **A**: **T1**
- **Initial operation** of data item **B**: **T1**
- **Final operation** of data item **A**: **T2**
- **Final operation** of data item **B**: **T2**
- Number of **RW** operations: **2**



S2

- **Initial operation** of data item **A**: **T1**
- **Initial operation** of data item **B**: **T1**
- **Final operation** of data item **A**: **T2**
- **Final operation** of data item **B**: **T2**
- Number of **RW** operations: **2**

SERIALIZABILITY

LOCKING

SERIALIZATION – GRAPH CHECKING

TIMESTAMP ORDERING

COMMITMENT ORDERING

MULTIVERSION CONCURRENCY CONTROL

INDEX CONCURRENCY CONTROL

PRIVATE WORKSPACE MODEL

LOCKING PROTOCOLS

SIMPLE LOCKING



BASIC 2PL

CONSERVATIVE 2PL

STRICT 2PL


RIGOROUS 2PL

SIMPLE LOCKING


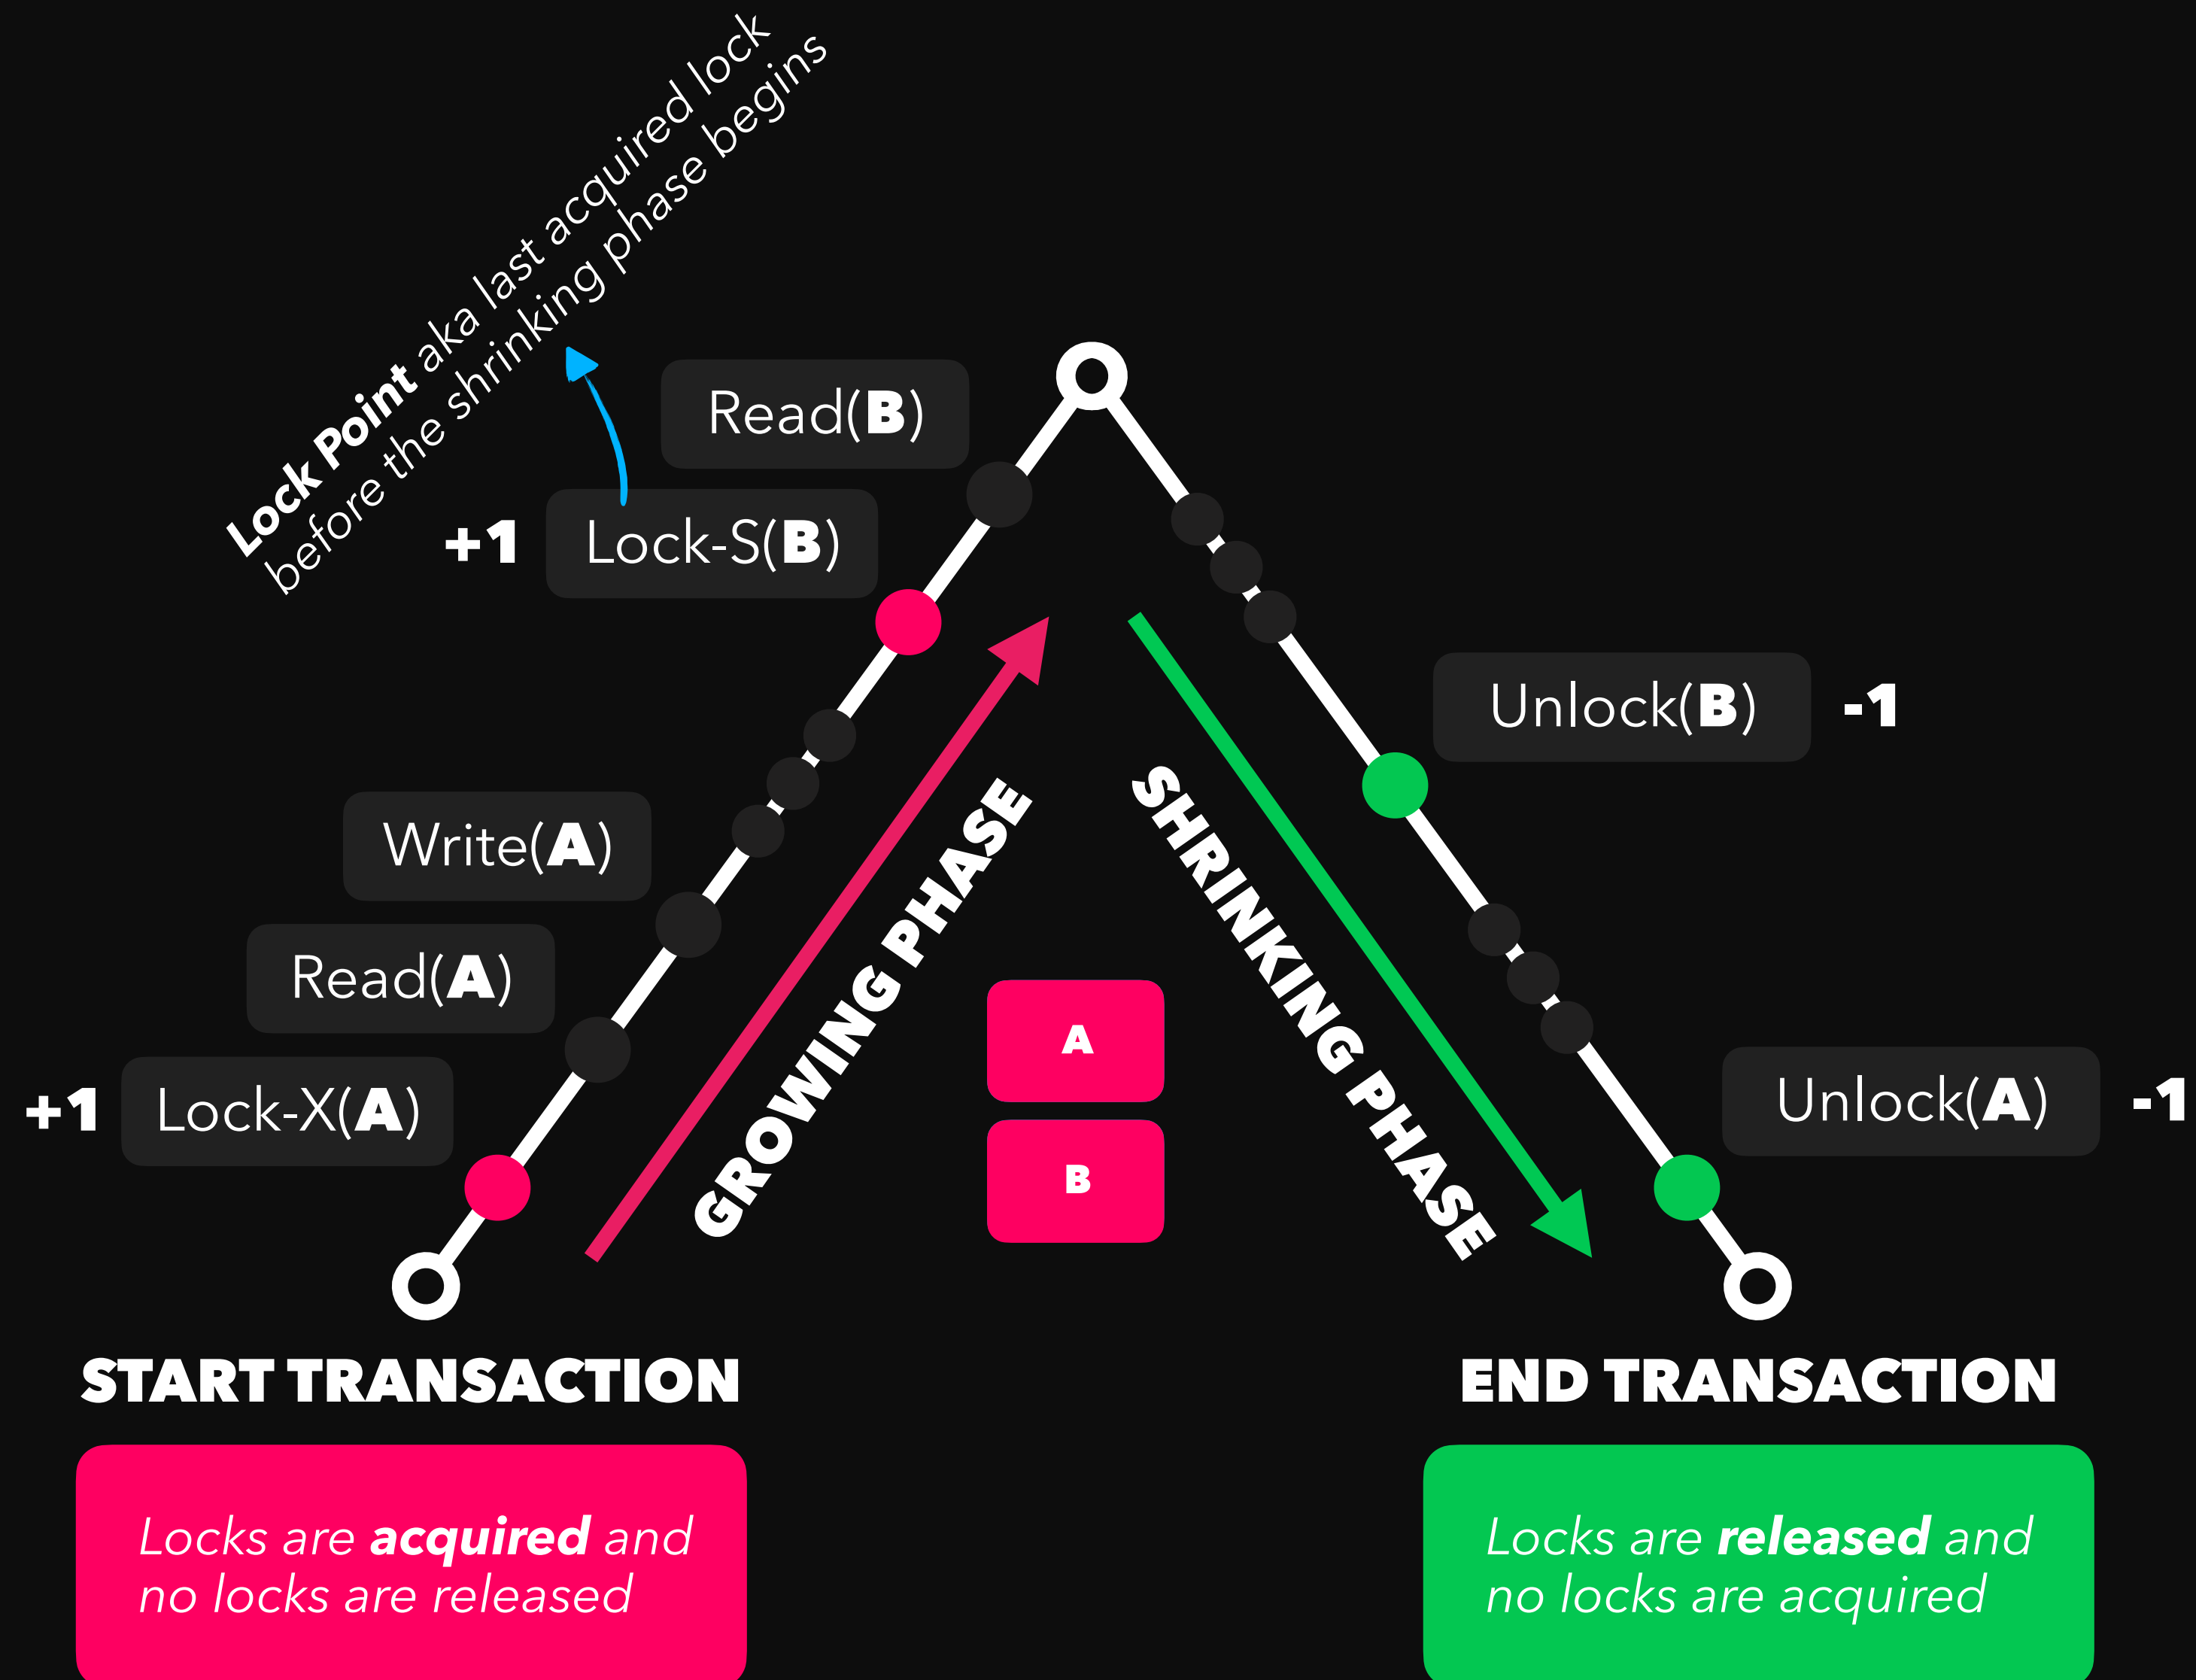


TYPE	READ LOCK	WRITE LOCK
READ LOCK	✓	X
WRITE LOCK	X	X

SIMPLE LOCKING EXAMPLE



BASIC 2PL

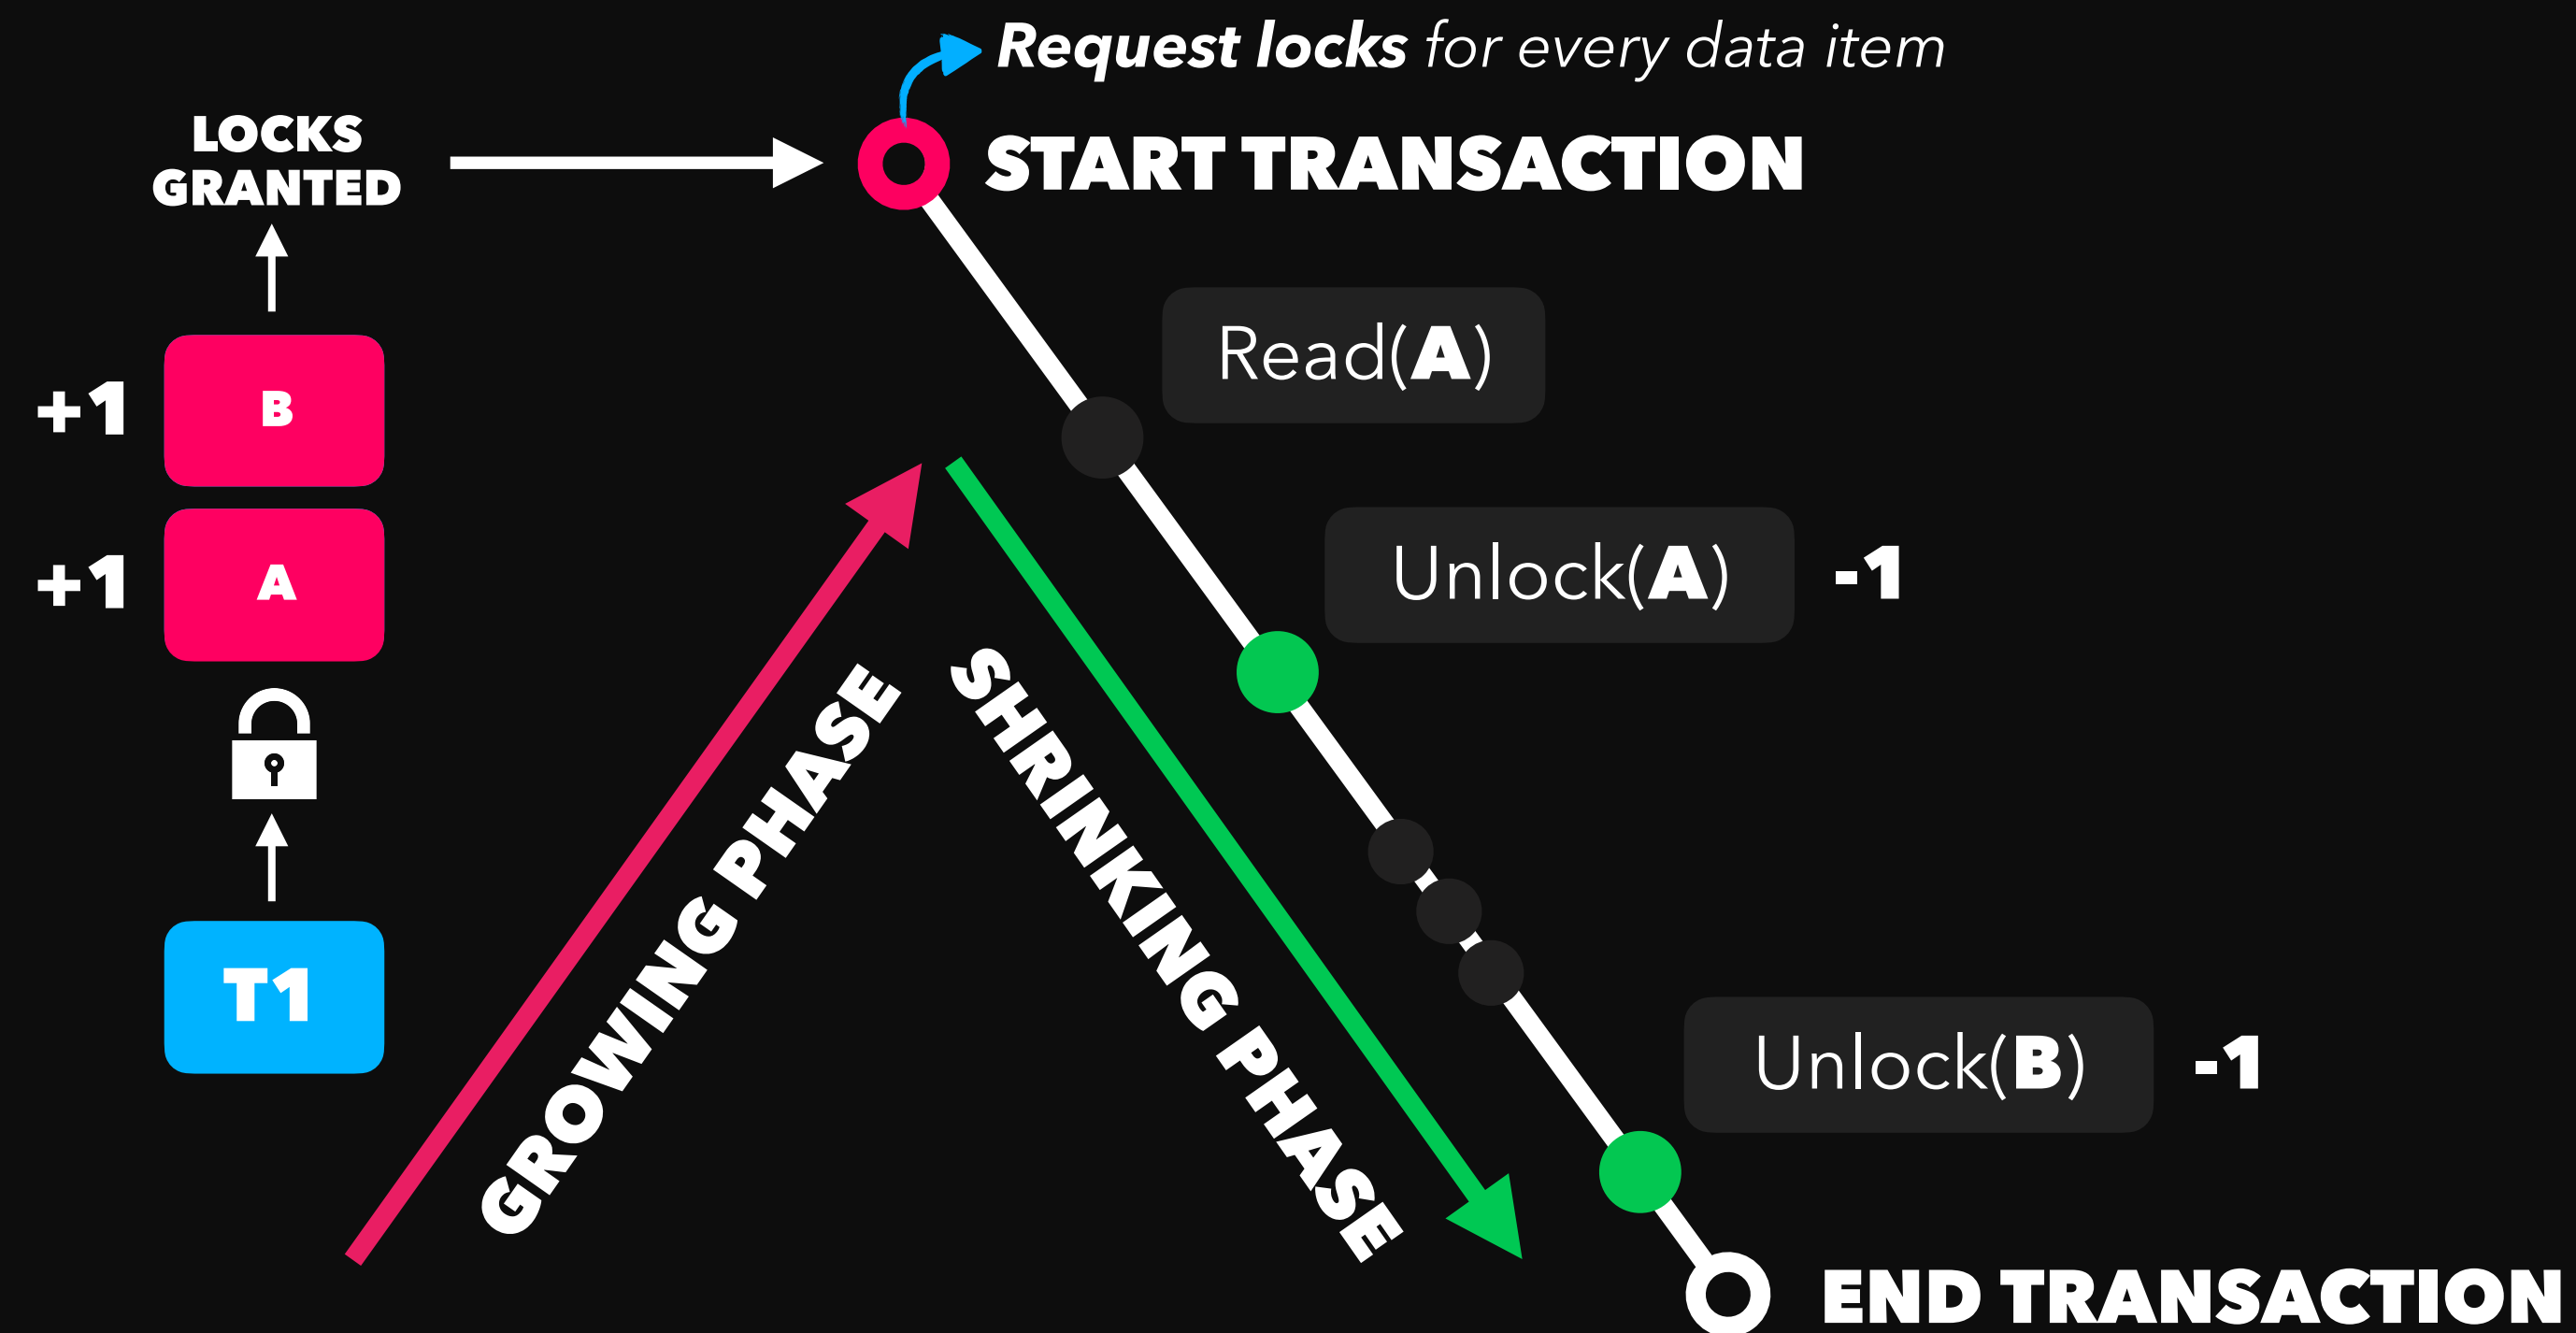


UNNECESSARY WAIT

DEADLOCKS

CASCADING ROLLBACKS

C2PL PROTOCOL



NO DEADLOCKS

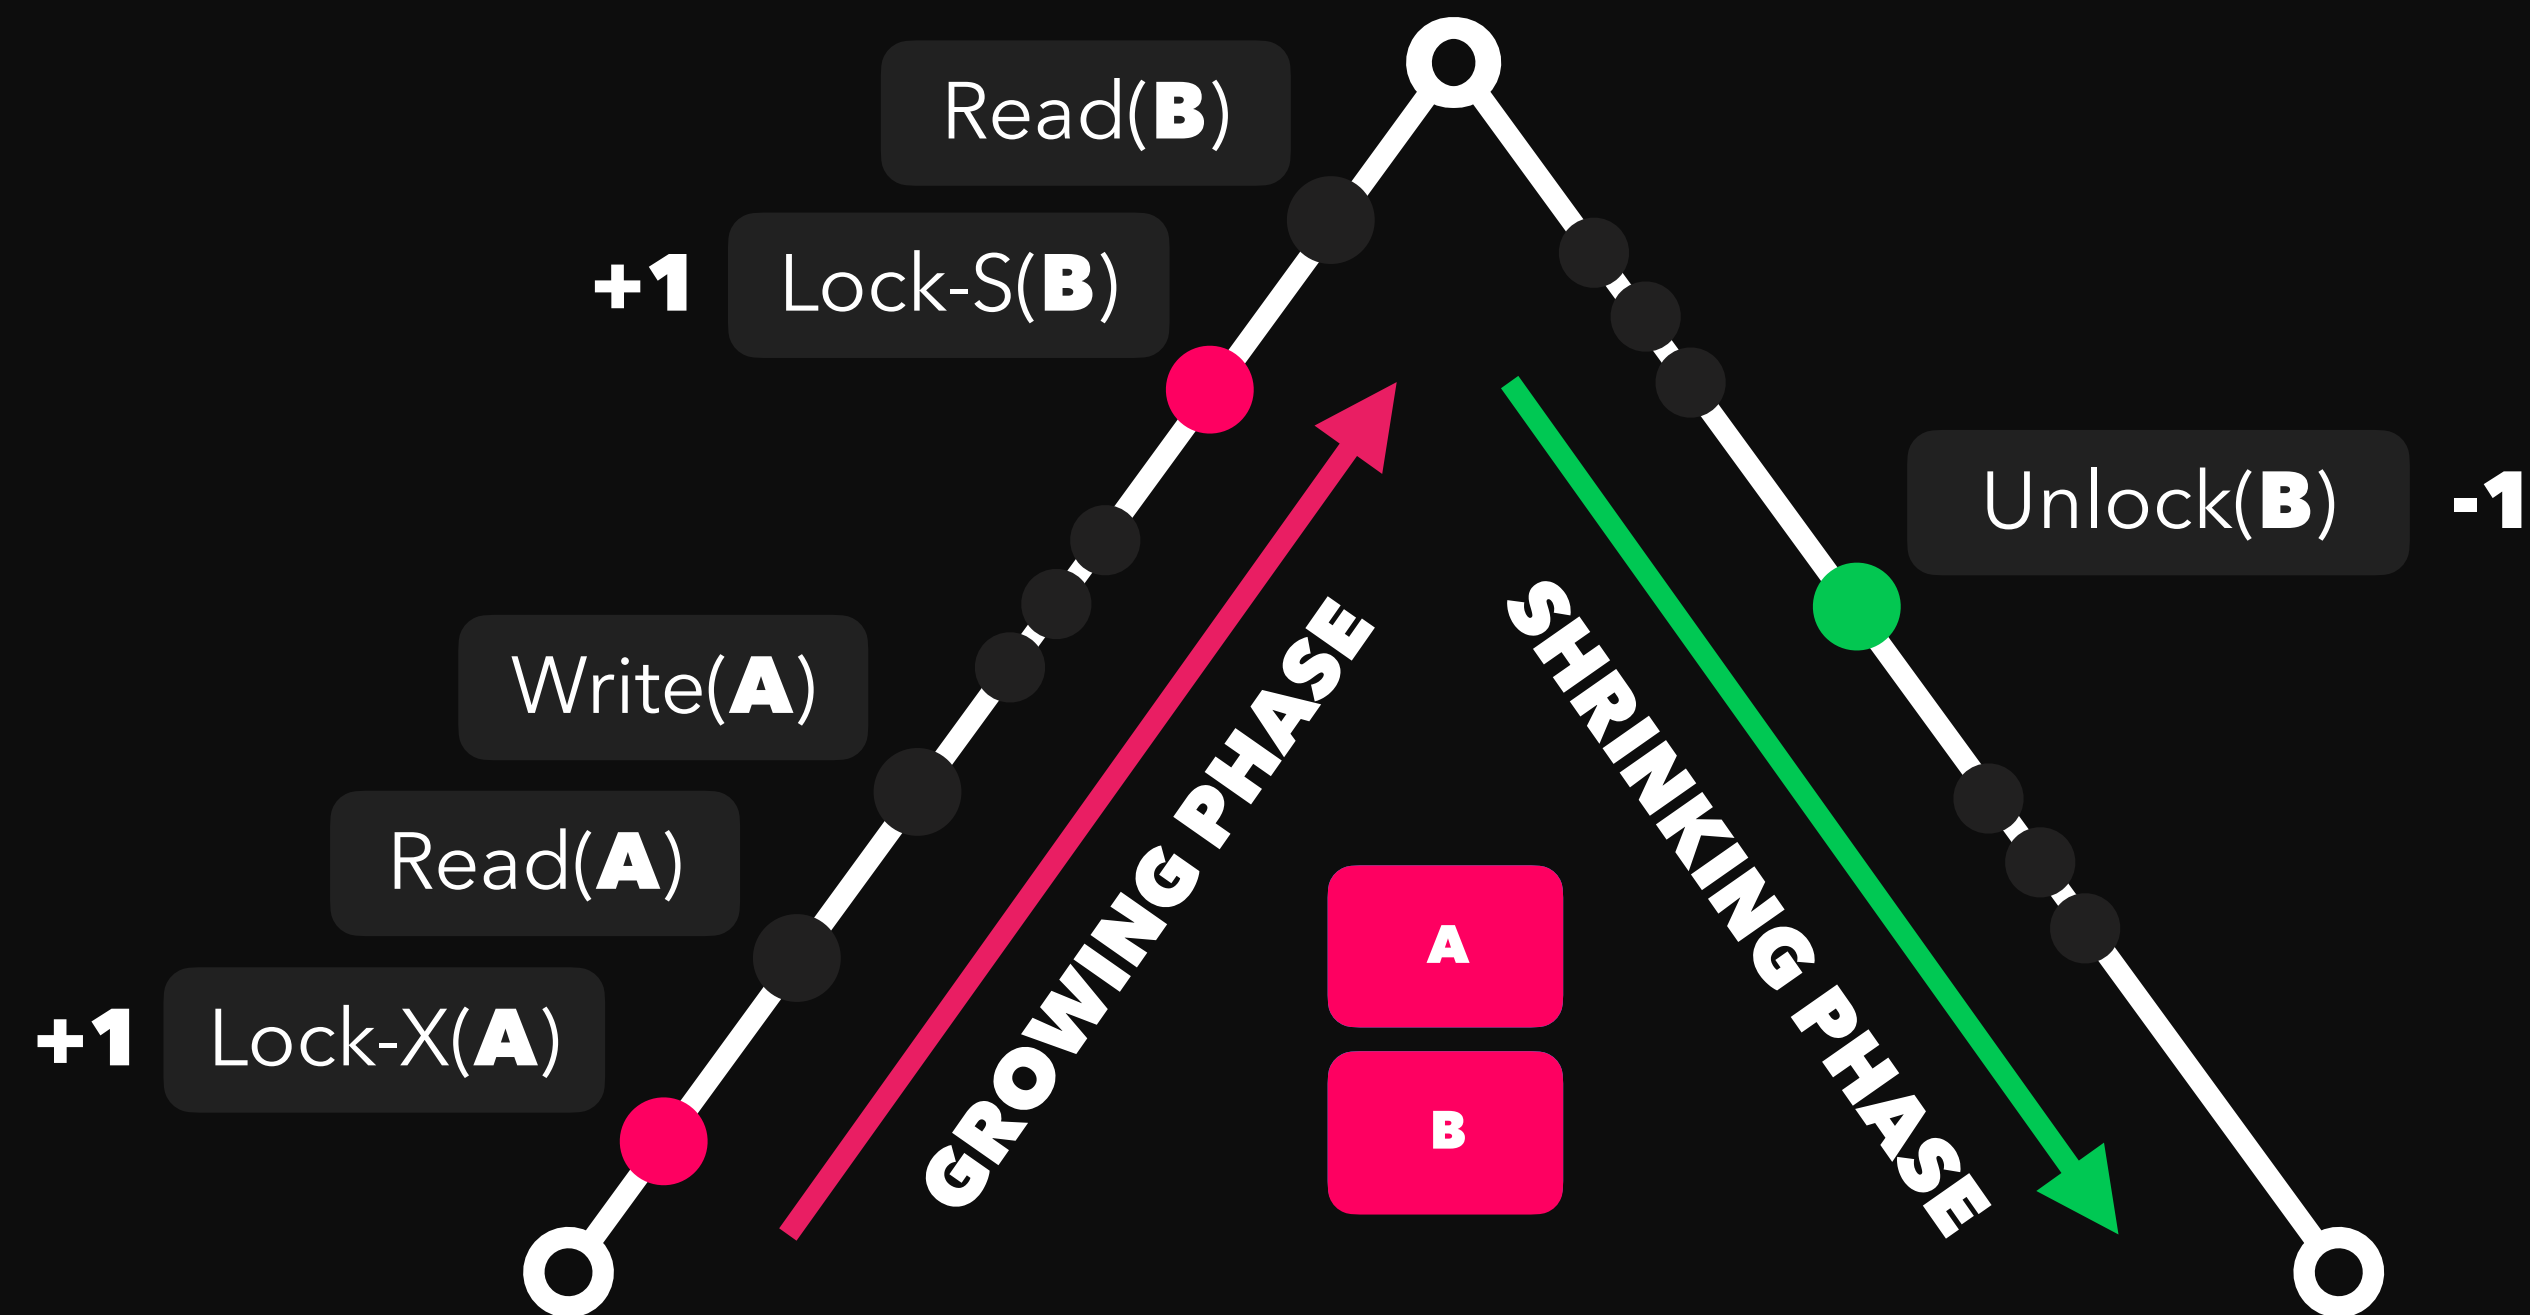
DIFICULT IMPLEMENTATION

CASCADING ROLLBACKS

*All locks are **acquired** before transaction **start***

*Locks are **released** and no locks are acquired*

S2PL PROTOCOL



START TRANSACTION

Locks are **acquired** and no locks are released

END TRANSACTION

Only **shared locks** are released

MOST POPULAR

STRICT SCHEDULING

EASY RECOVERY

NO CASCADING ROLLBACKS

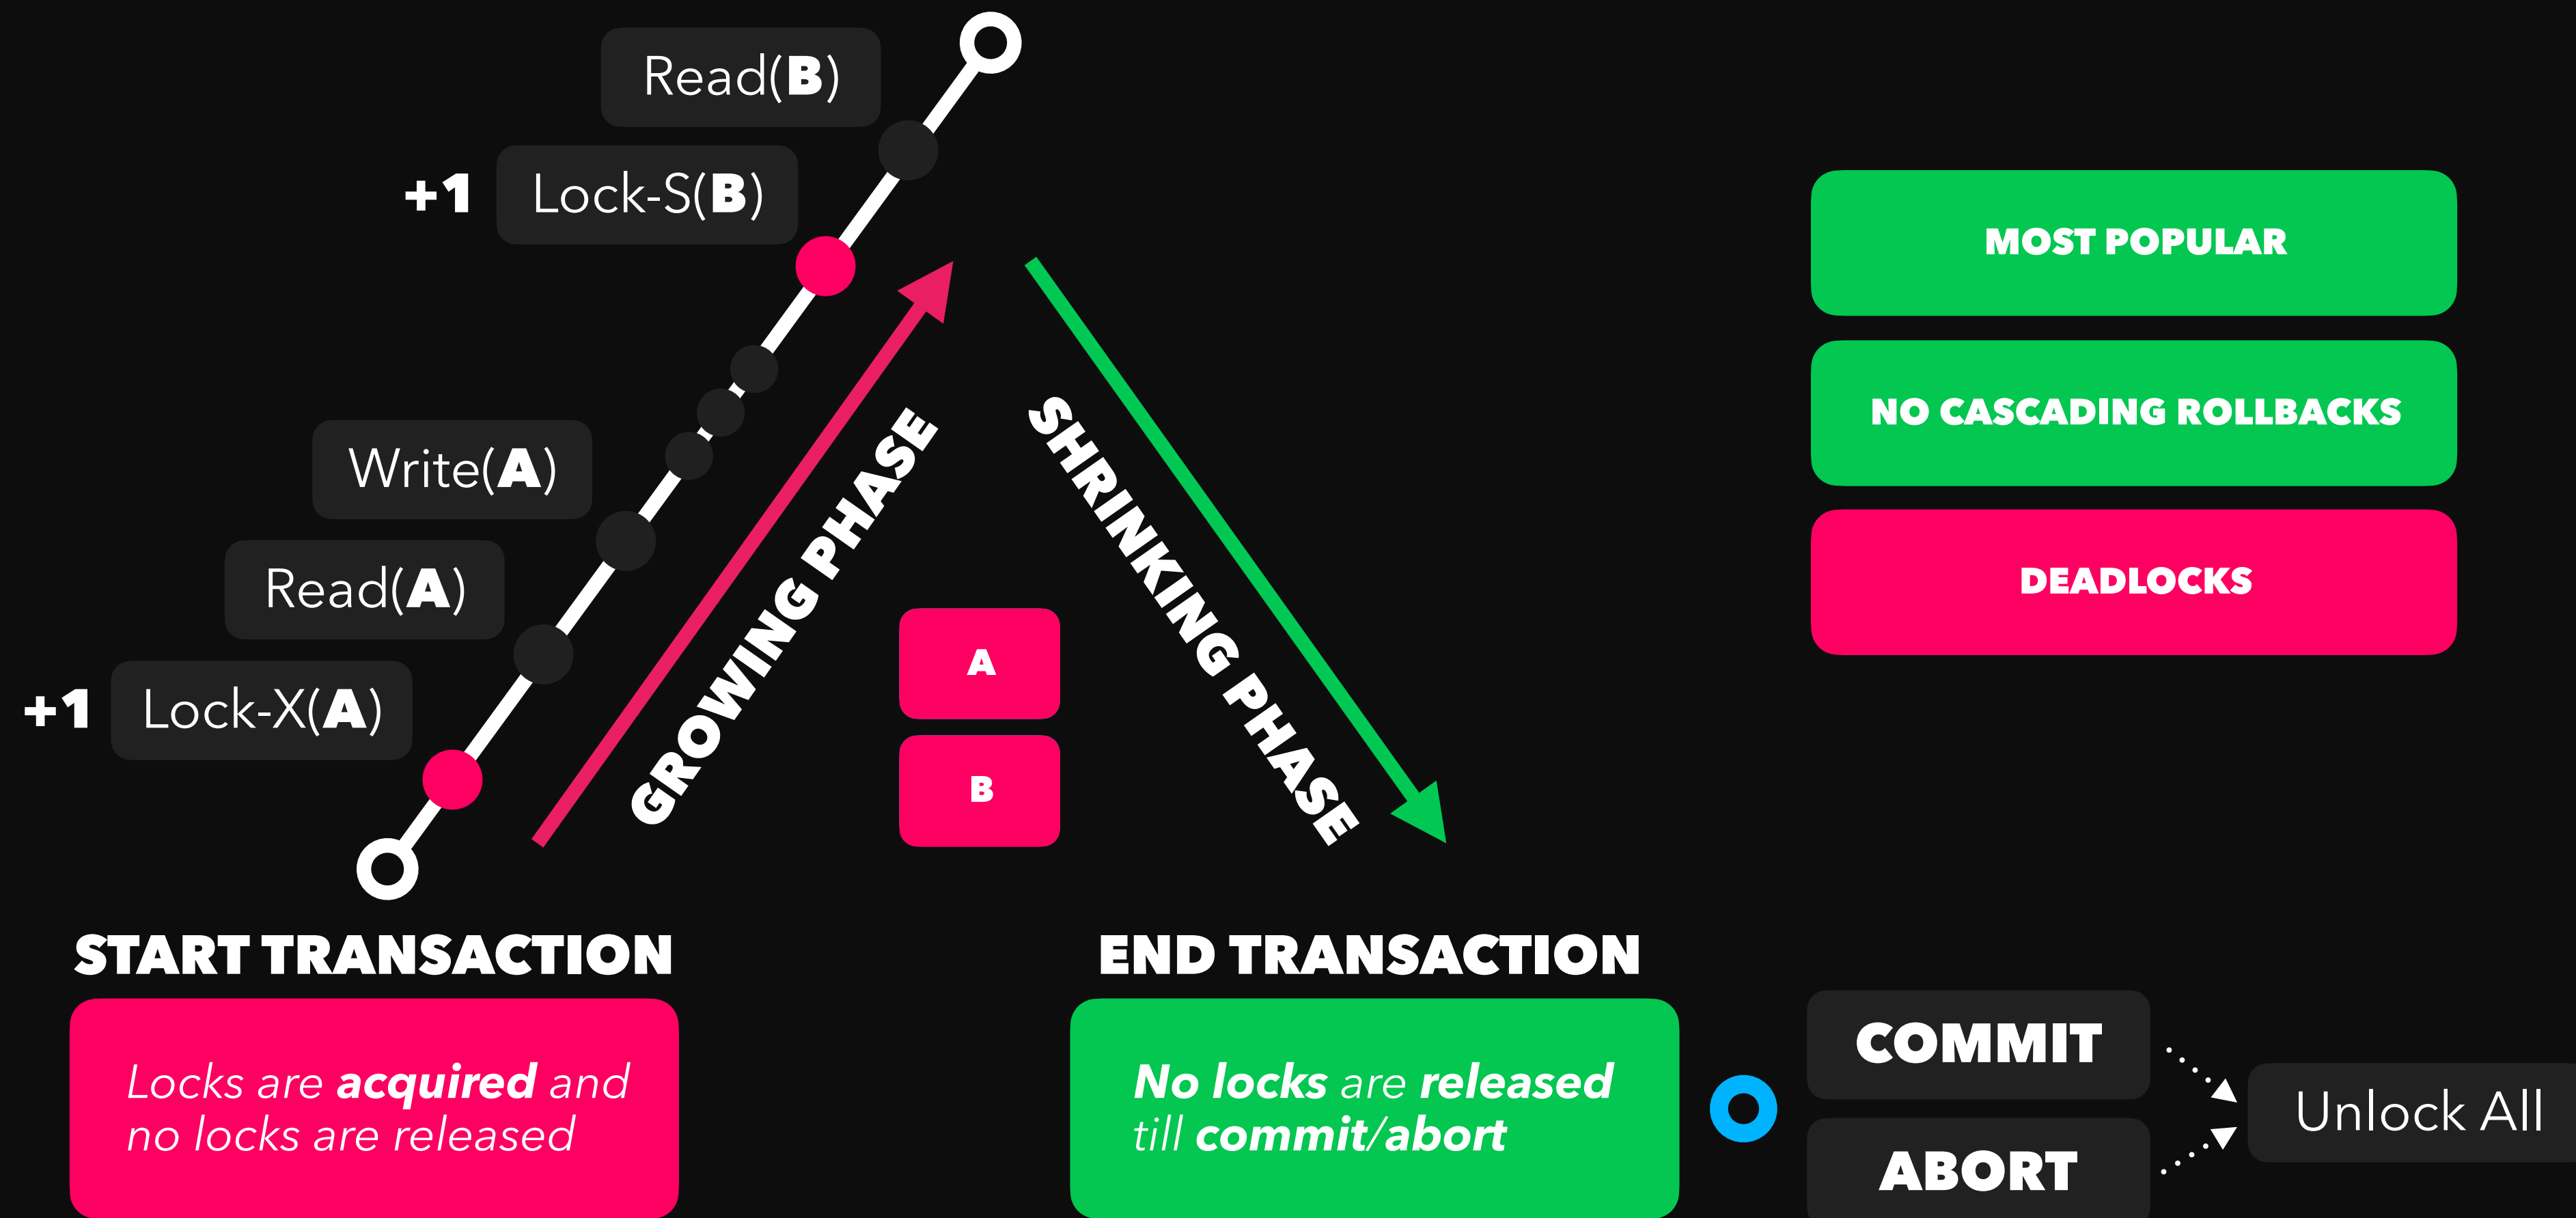
DEADLOCKS

COMMIT

ABORT


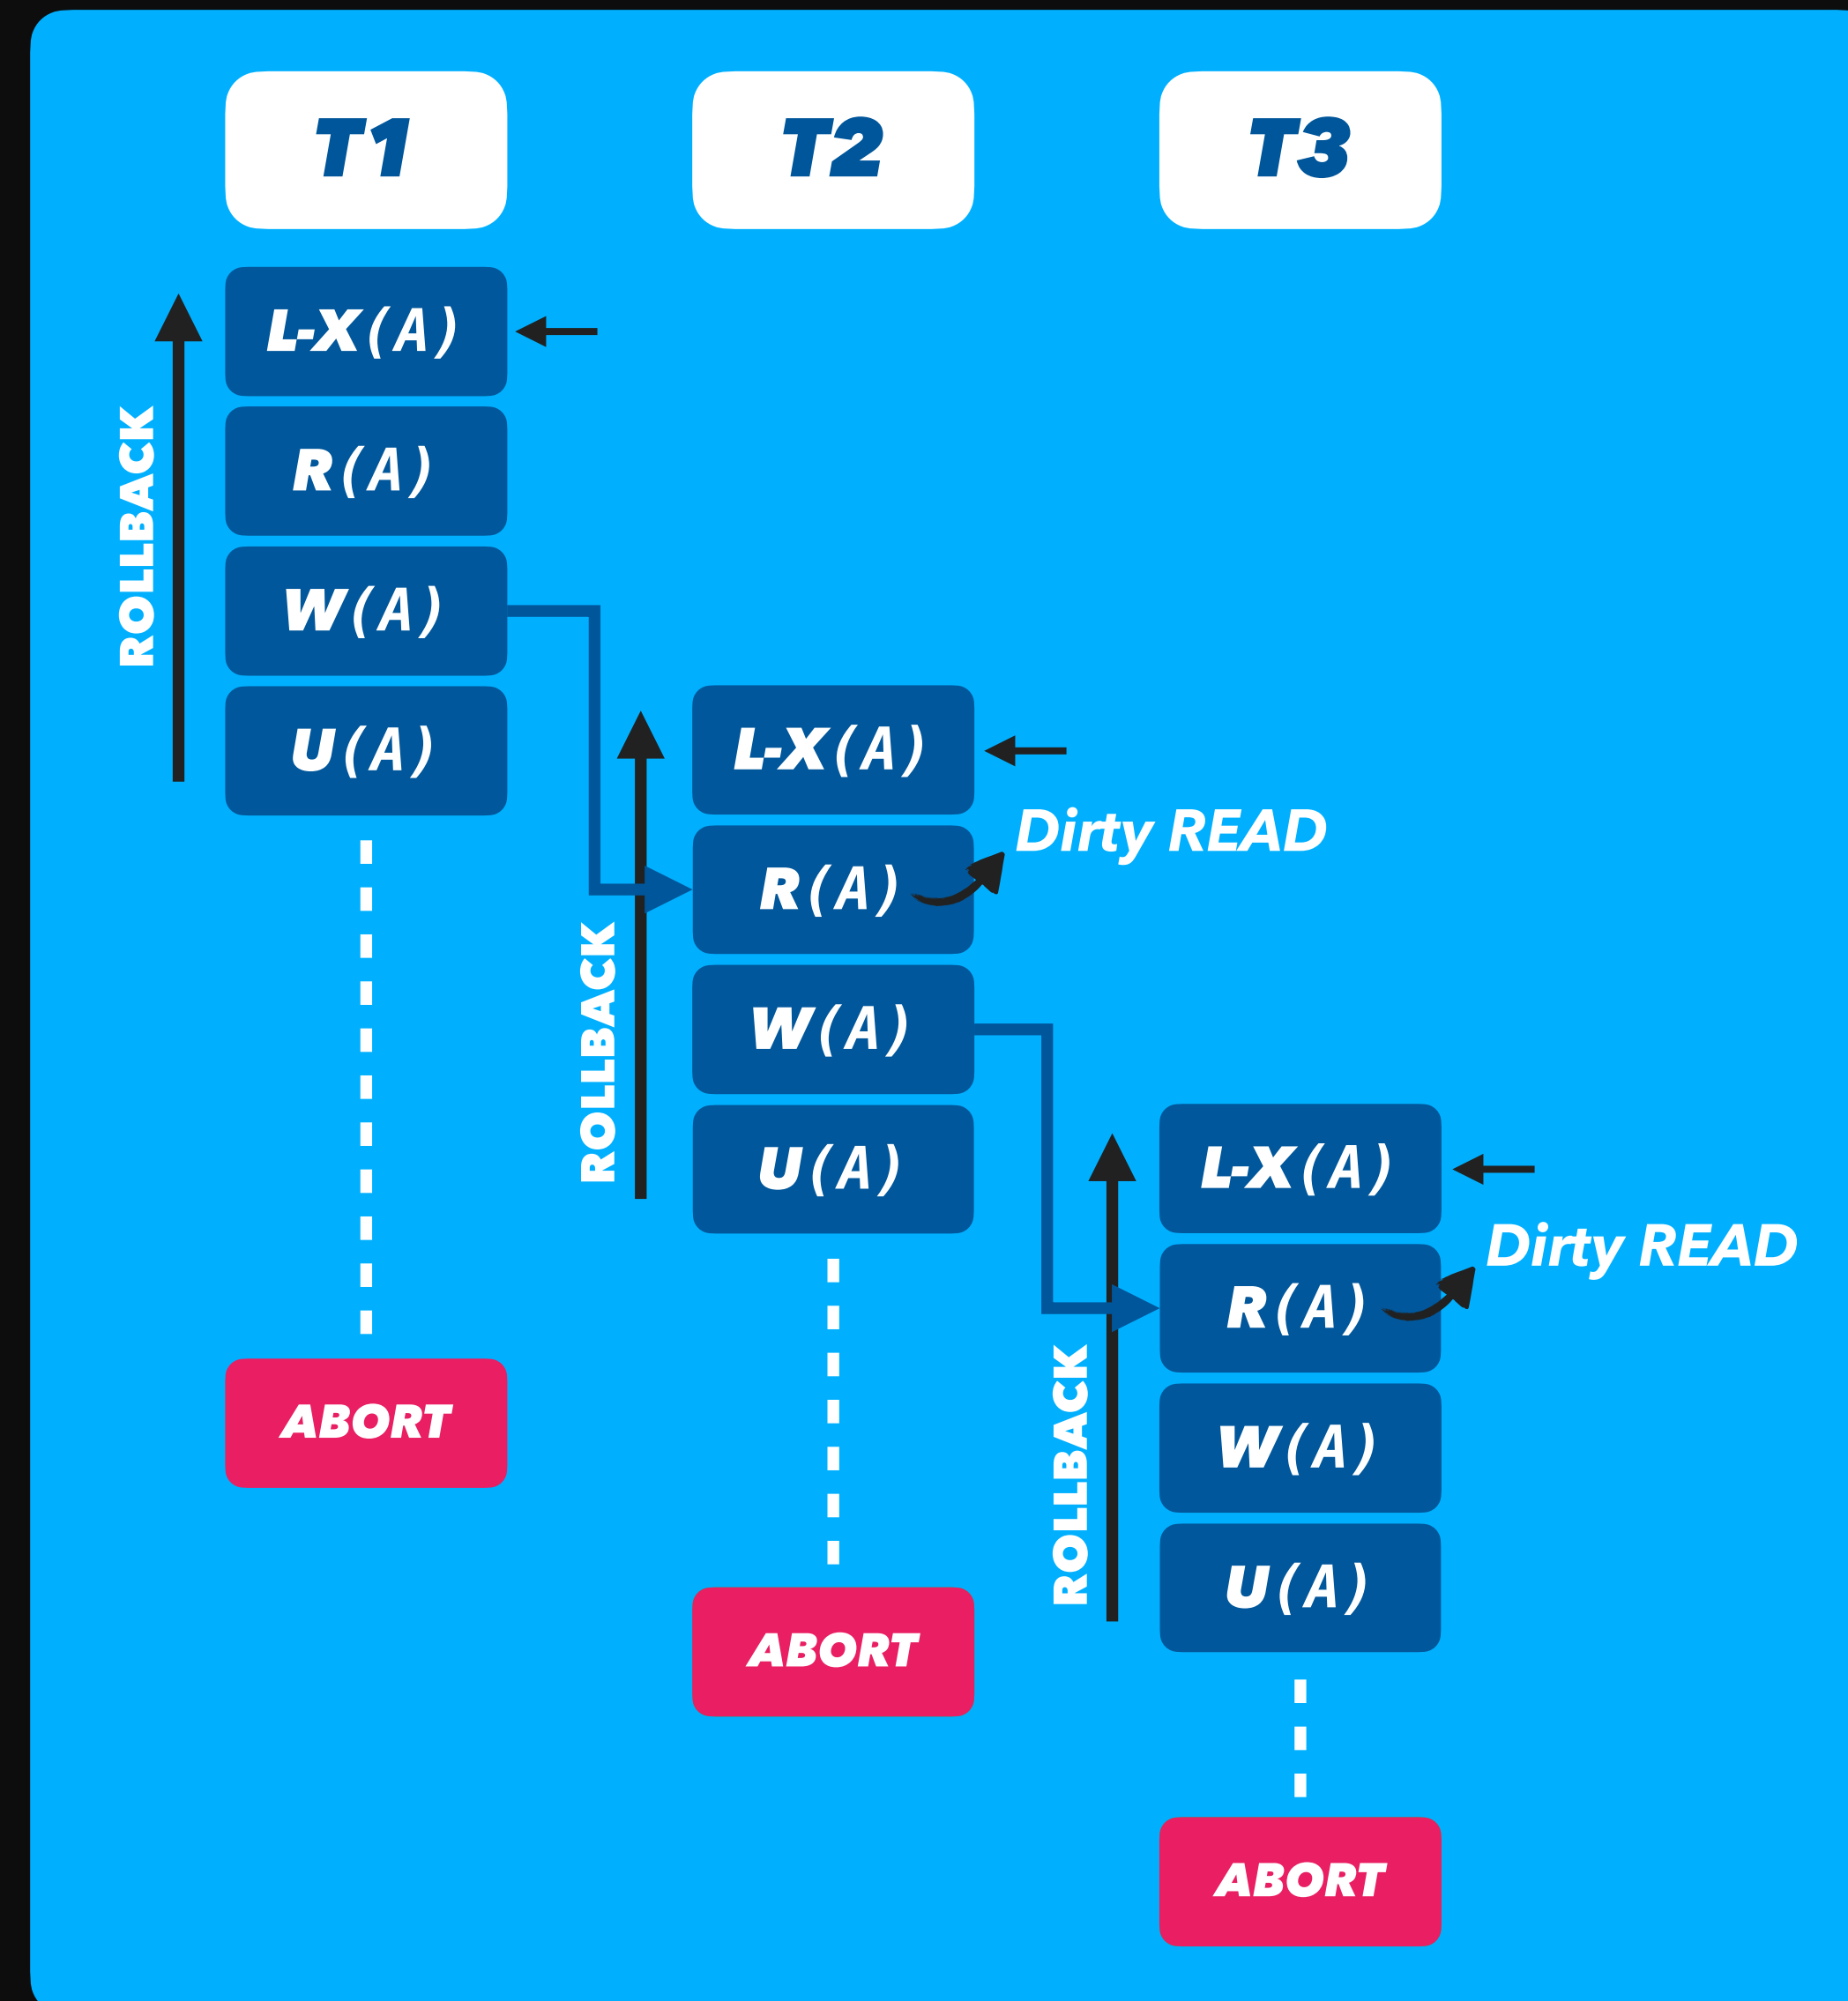
Unlock(A) -1

SS2PL/R2PL PROTOCOL



RECOVERABILITY


**CASCADING
ABORT**



TRANSACTION ISSUES

LOST UPDATE

DIRTY READ


INCORRECT SUMMARY

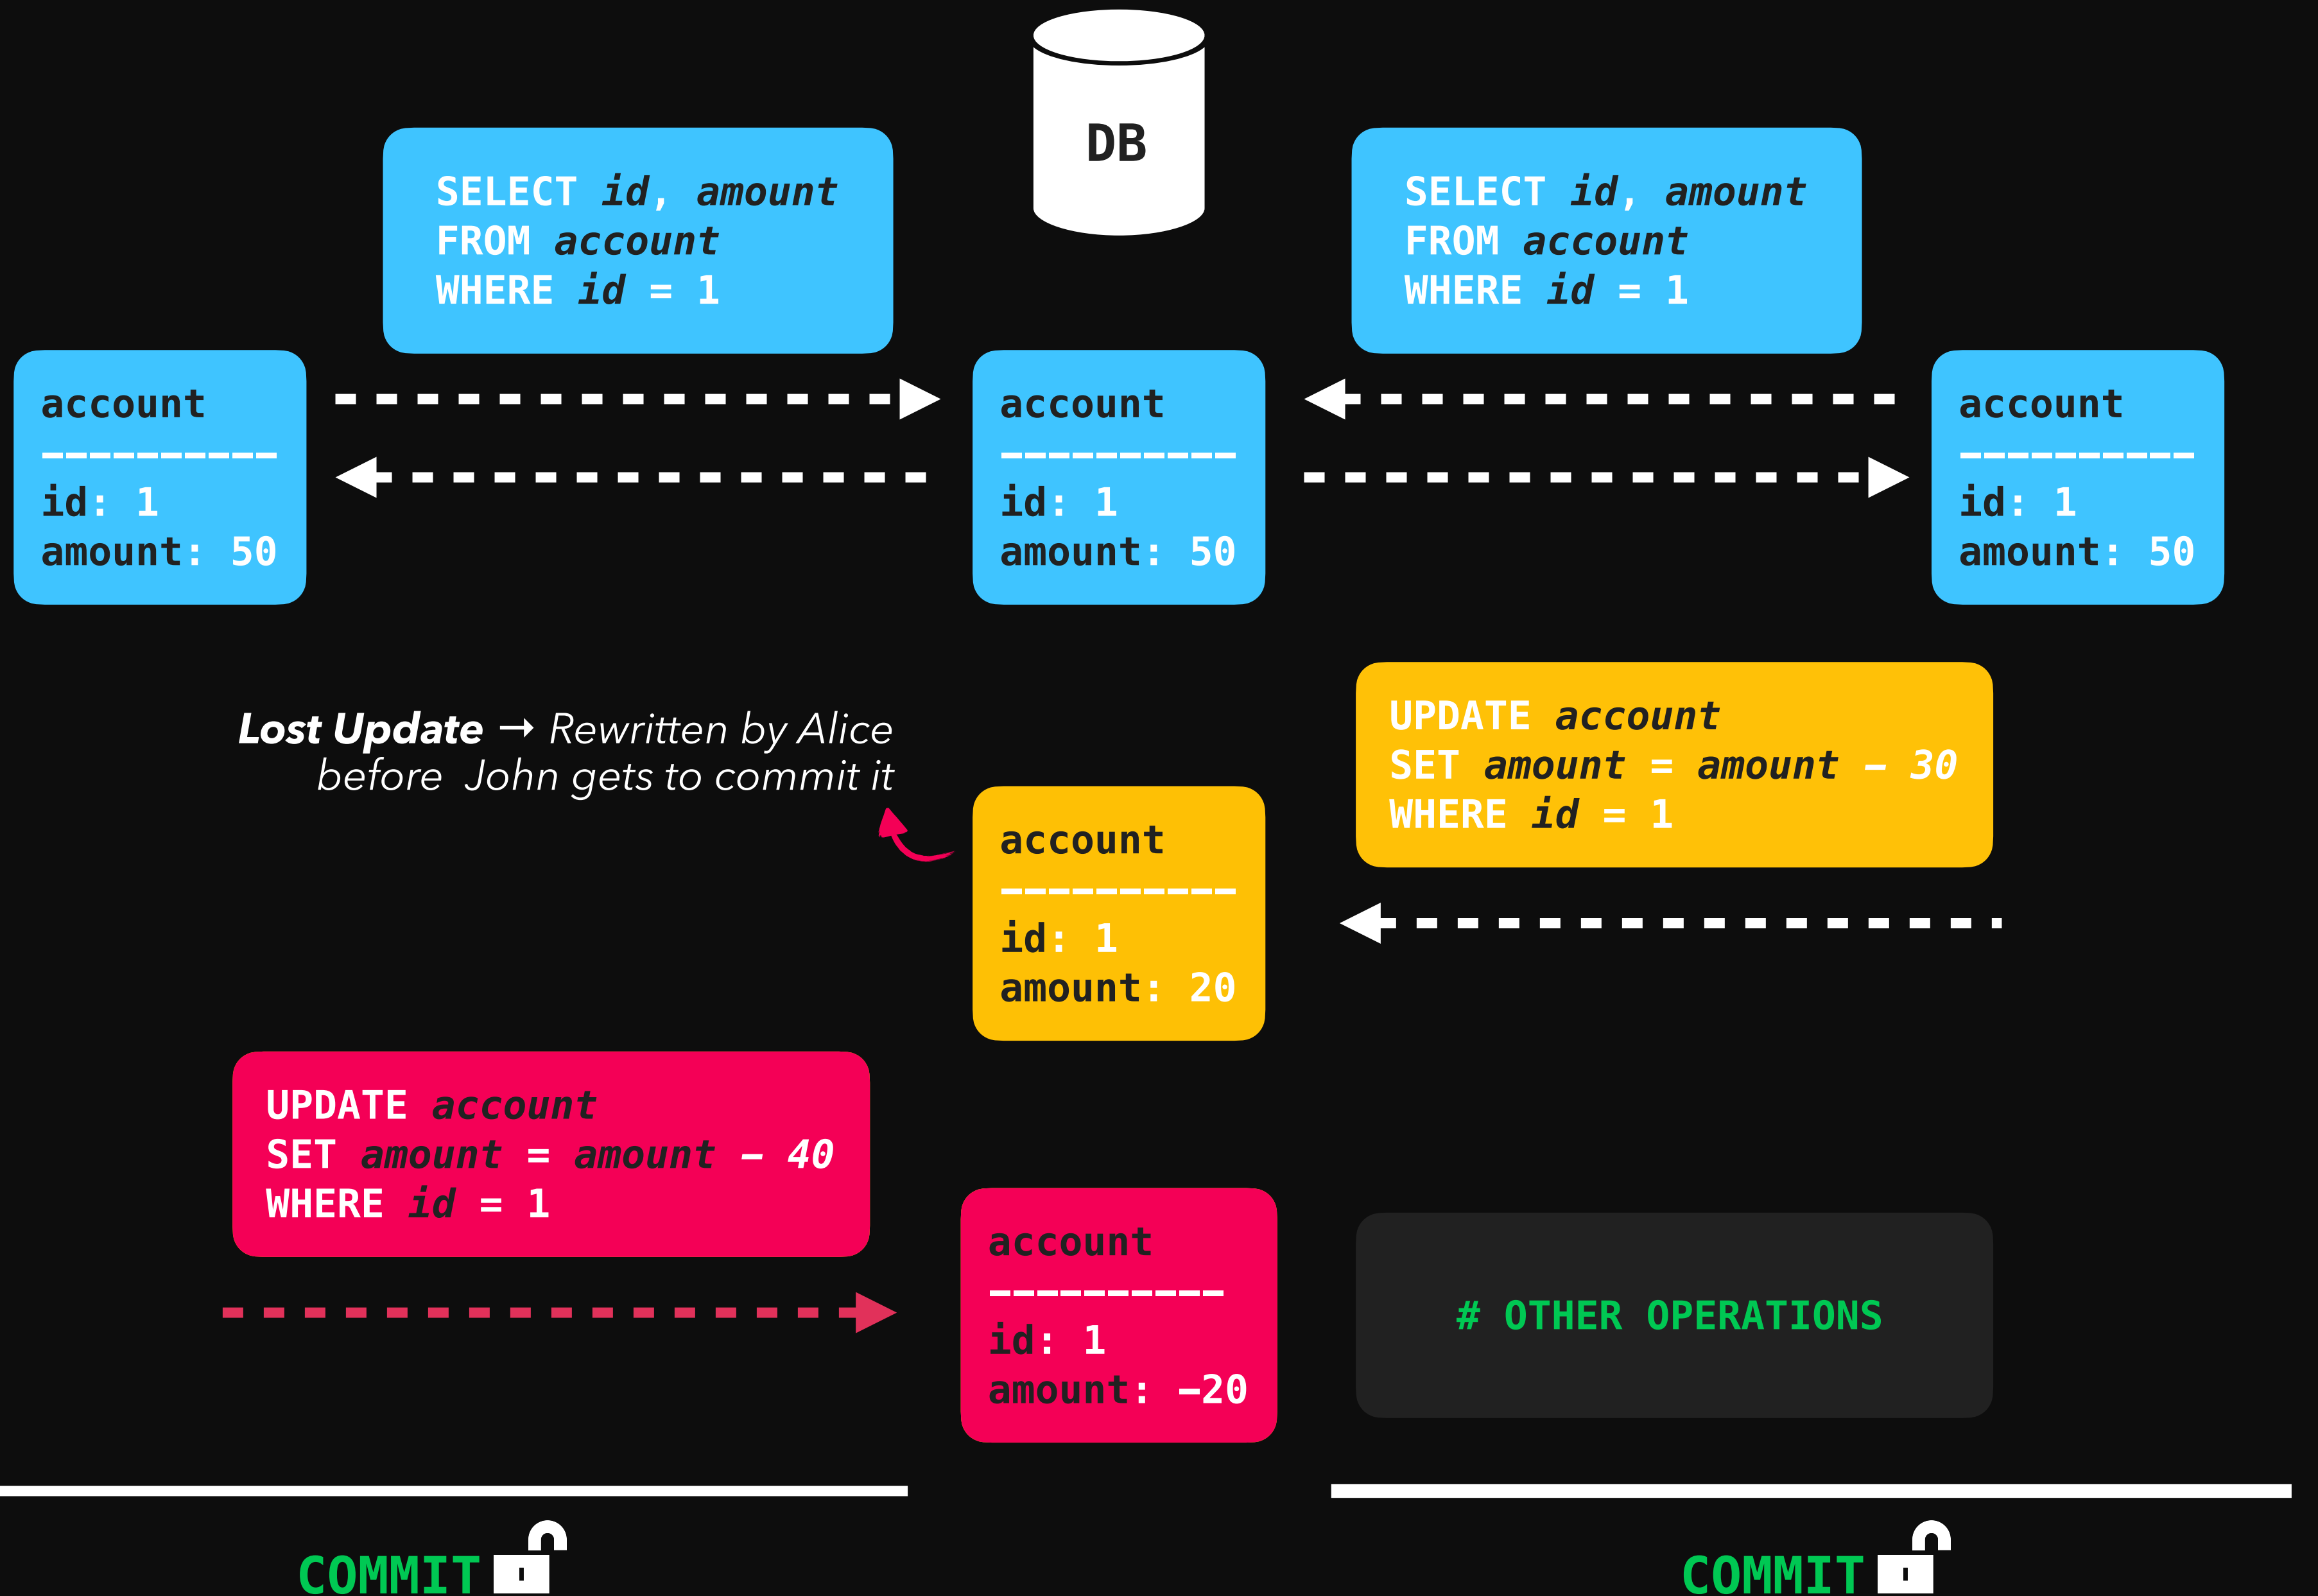
LOCKING TYPES

OPTIMISTIC

PESSIMISTIC

SEMI-OPTIMISTIC

T1
↑

Alice



T2
↑

John



Alice

account

id: 1
amount: 50

OTHER OPERATIONS

COMMIT 

SELECT *id, amount*
FROM *account*
WHERE *id* = 1

Shared Lock



Exclusive Lock

SELECT *id, amount*
FROM *account*
WHERE *id* = 1

account

id: 1
amount: 50

account

id: 1
amount: 20

UPDATE *account*
SET *amount* = *amount* - 30
WHERE *id* = 1

COMMIT 

account

id: 1
amount: 50



John



Alice

